



PES UNIVERSITY

Department of Computer Science & Engineering

Computer Network Security

UE23CS343AB6

Assignment 7 Submission

Name of the Student	Pranav Hemanth
SRN	PES1UG23CS433
Section	G
Department	CSE
Campus	RR

Computer Network Security

UE23CS343AB6

Explanation of the config files as part of this lab setup:

This lab uses a small multi-container environment (router container + helper code) implemented with Docker images, source code and host-mounted volumes. The top-level tree is:

```
[pranavhemanth@Pranavs-MacBook-Pro-M3 Lab7 %ls
 docker-compose.yml      router
 Files                  volumes
[pranavhemanth@Pranavs-MacBook-Pro-M3 Lab7 %tree
 .
├── docker-compose.yml
└── Files
    ├── kernel_module
    │   └── hello.c
    │       └── Makefile
    └── packet_filter
        ├── Makefile
        └── seedFilter.c
└── router
    └── Dockerfile
volumes
    ├── kernel_module
    │   └── hello.c
    │       └── Makefile
    └── packet_filter
        ├── Makefile
        ├── seedBlock.c
        ├── seedFilter.c
        └── seedPrint.c

8 directories, 12 files
pranavhemanth@Pranavs-MacBook-Pro-M3 Lab7 %
```

Below is the role and purpose of each file and its configurations:

docker-compose.yml

Defines the multi-container composition for the lab. It describes services to bring up (for example, a router service), how volumes are mounted from ./volumes into containers, network configuration, and any runtime privileges (e.g., privileged: true or cap_add: NET_ADMIN) required to load kernel modules or use raw sockets. This is the entrypoint for launching the lab environment with a single docker-compose up command.

Files/

A repository copy of the source used in the lab. These are kept as reference or for building locally before being mounted into containers.

- Files/kernel_module/hello.c
A simple example kernel module source (commonly a “hello world” kernel module). Used to demonstrate how to compile, load, and unload a kernel module inside a privileged container or on the host during the lab.
- Files/kernel_module/Makefile
Makefile to compile the kernel module (hello.c) against the appropriate kernel headers. Shows compile targets such as make, make clean, and possibly an install/insmod helper.
- Files/packet_filter/seedFilter.c
A user-space packet filtering program (single source) that demonstrates packet capture/filtering logic or integration points with the kernel module. The Makefile compiles this into a runnable binary used to observe/filter traffic inside the router container.

router/

Contains the Dockerfile used to build the router container image that runs the lab’s networking components (packet filter, kernel module build tools, or test scripts).

- router/Dockerfile
The image definition used for the router service. Typical responsibilities:
 - Starts from a base image (e.g., ubuntu), installs required build tools (build-essential, linux-headers), and network utilities (ip, tcpdump) so you can compile and load kernel modules and run packet filtering binaries inside the container.
 - Copies or expects mounted volumes (from volumes/) so you can build from the host-provided source.

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

- May set an entrypoint or default command to keep the container running (e.g., a small shell wrapper or supervisor) so you can docker exec in to build/load modules or run tests.

volumes/

Host-side directories that are mounted into containers by docker-compose.yml. These contain the actual code the containers run or build. Using volumes makes iteration easy: edit files on your host and the container sees the changes immediately.

- volumes/kernel_module/hello.c
Same kernel module source as in Files/ but placed here to be mounted into the router container. You typically docker exec into the container, cd /mnt/volumes/kernel_module, then make and insmod ./hello.ko (requires privileged container and matching kernel headers).
- volumes/kernel_module/Makefile
The build instructions for the kernel module used inside the container. Keeps compilation consistent when building inside the lab container.
- volumes/packet_filter/Makefile
Makefile to build the packet filter and other small user-space utilities. When mounted into the router container you run make to produce the binaries.
- volumes/packet_filter/seedBlock.c
A source file that likely implements blocking logic (e.g., dropping certain packets). Used as one behavior of the packet-filtering toolset.
- volumes/packet_filter/seedFilter.c
A packet capture / filter implementation — probably the main filtering binary that inspects packets and optionally hands decisions to seedBlock or logs decisions.
- volumes/packet_filter/seedPrint.c
A helper program that prints/logs packets or events (useful to observe what the filter sees). Helpful for debugging and lab demonstrations (shows matched packets, counters, etc).

Change shell name:

PS1="terminal1:PES1UG23CS433:PranavHemanth:\w\n\\$>"

PS1="terminal2:PES1UG23CS433:PranavHemanth:\w\n\\$>"

PS1="hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:\w\n\\$>"

PS1="seed-router:PES1UG23CS433:PranavHemanth:\w\n\\$>"

PS1="host1-192.168.60.5:PES1UG23CS433:PranavHemanth:\w\n\\$>"

PS1="host2-192.168.60.6:PES1UG23CS433:PranavHemanth:\w\n\\$>"

PS1="host3-192.168.60.7:PES1UG23CS433:PranavHemanth:\w\n\\$>"

Notes about containers-

Since all the containers share the same kernel, kernel modules are global.

Therefore, if we set a kernel module from a container, it affects all the containers and the host. For this reason, it does not matter where you set the kernel module. We will just set the kernel module from the host VM in this lab.

Another thing to keep in mind is that containers' IP addresses are virtual. Packets going to these virtual IP addresses may not traverse the same path as what is described in the Netfilter document.

Therefore, in this task, to avoid confusion, we will try to avoid using those virtual addresses. We do most tasks on the host VM. The containers are mainly for other tasks.

Task 1: Implementing a Simple Firewall

In this task, we will implement a simple packet filtering type of firewall, which inspects each incoming and outgoing packet, and enforces the firewall policies set by the administrator.

Task 1.A: Implement a Simple Kernel Module

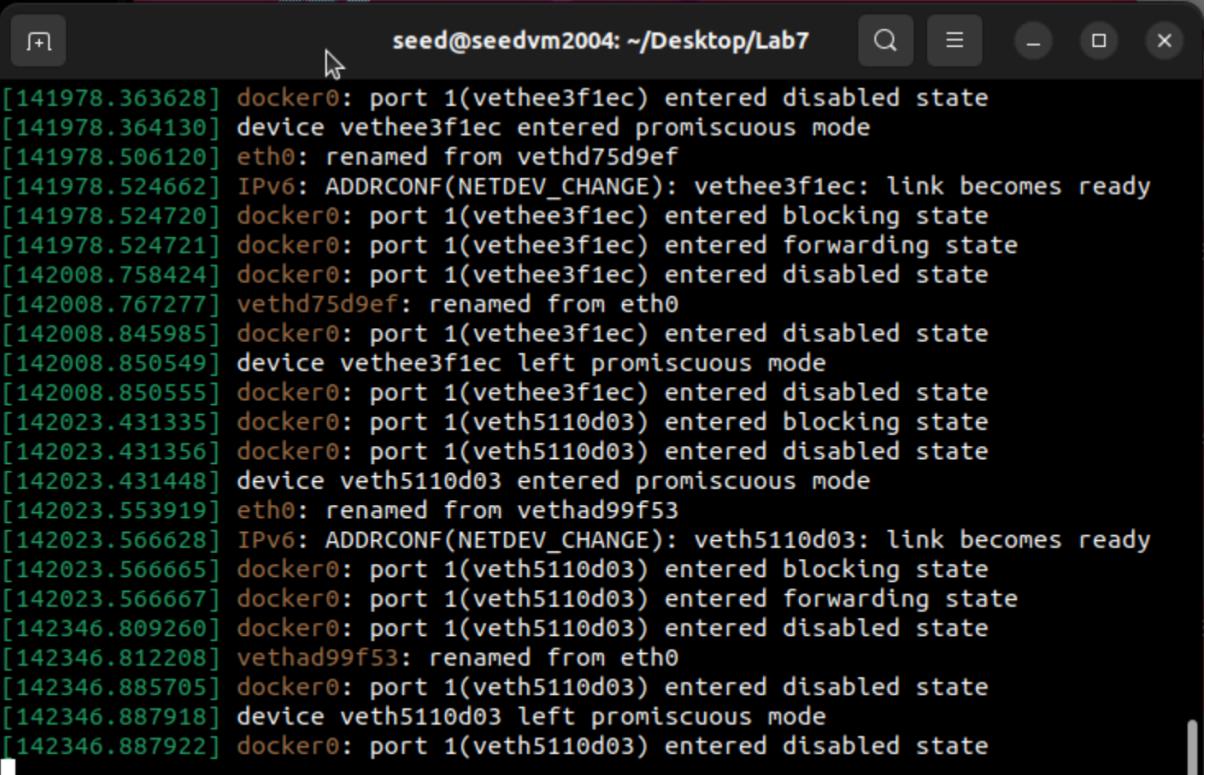
LKM allows us to add a new module to the kernel at runtime. This new module enables us to extend the functionalities of the kernel, without rebuilding the kernel or even rebooting the computer. The packet filtering part of a firewall can be implemented as an LKM.

Step1: On the VM- Open two Terminal Tabs, one to load the module and the other to view the messages.

terminal 1:

- 1) Initial Kernel Buffer:

Command: sudo dmesg -k -w



A screenshot of a terminal window titled "seed@seedvm2004: ~/Desktop/Lab7". The window displays a continuous stream of kernel log messages. The messages are timestamped and show various network interface states (disabled, promiscuous mode, blocking, forwarding) for interfaces like docker0, veth, and eth. Some entries indicate IPv6 ADDRCONF events. The log ends with a message from the user "PranavHemanth" at the bottom.

```
[141978.363628] docker0: port 1(vethee3f1ec) entered disabled state
[141978.364130] device vethee3f1ec entered promiscuous mode
[141978.506120] eth0: renamed from vethd75d9ef
[141978.524662] IPv6: ADDRCONF(NETDEV_CHANGE): vethee3f1ec: link becomes ready
[141978.524720] docker0: port 1(vethee3f1ec) entered blocking state
[141978.524721] docker0: port 1(vethee3f1ec) entered forwarding state
[142008.758424] docker0: port 1(vethee3f1ec) entered disabled state
[142008.767277] vethd75d9ef: renamed from eth0
[142008.845985] docker0: port 1(vethee3f1ec) entered disabled state
[142008.850549] device vethee3f1ec left promiscuous mode
[142008.850555] docker0: port 1(vethee3f1ec) entered disabled state
[142023.431335] docker0: port 1(veth5110d03) entered blocking state
[142023.431356] docker0: port 1(veth5110d03) entered disabled state
[142023.431448] device veth5110d03 entered promiscuous mode
[142023.553919] eth0: renamed from vethad99f53
[142023.566628] IPv6: ADDRCONF(NETDEV_CHANGE): veth5110d03: link becomes ready
[142023.566665] docker0: port 1(veth5110d03) entered blocking state
[142023.566667] docker0: port 1(veth5110d03) entered forwarding state
[142346.809260] docker0: port 1(veth5110d03) entered disabled state
[142346.812208] vethad99f53: renamed from eth0
[142346.885705] docker0: port 1(veth5110d03) entered disabled state
[142346.887918] device veth5110d03 left promiscuous mode
[142346.887922] docker0: port 1(veth5110d03) entered disabled state
```

terminal 2:

2) Commands:

```
$ make
```

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/kernel_module
$make
make -C /lib/modules/5.15.0-153-generic/build M=/home/seed/Desktop/Lab7/volumes/
kernel_module modules
make[1]: Entering directory '/usr/src/linux-headers-5.15.0-153-generic'
warning: the compiler differs from the one used to build the kernel
  The kernel was built by: gcc (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0
  You are using:          gcc (Ubuntu 11.4.0-1ubuntu1~22.04.2) 11.4.0
make[1]: Leaving directory '/usr/src/linux-headers-5.15.0-153-generic'
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/kernel_module
$
```

```
$ sudo insmod hello.ko (inserting a module)
```

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/kernel_module
$ls
hello.c  hello.mod  hello.mod.o  Makefile      Module.symvers
hello.ko  hello.mod.c  hello.o       modules.order
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/kernel_module
$sudo insmod hello.ko
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/kernel_module
$
```

```
$ lsmod | grep hello (list modules)
```

```
$lsmod | grep hello
hello           16384  0
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/kernel_module
$
```

```
$ sudo rmmod hello
```

```
$sudo rmmod hello  
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/kernel_module  
$
```

terminal 1:

```
[141978.524721] docker0: port 1(veth3f1ec) entered forwarding state  
[142008.758424] docker0: port 1(veth3f1ec) entered disabled state  
[142008.767277] vethd75d9ef: renamed from eth0  
[142008.845985] docker0: port 1(veth3f1ec) entered disabled state  
[142008.850549] device veth3f1ec left promiscuous mode  
[142008.850555] docker0: port 1(veth3f1ec) entered disabled state  
[142023.431335] docker0: port 1(veth5110d03) entered blocking state  
[142023.431356] docker0: port 1(veth5110d03) entered disabled state  
[142023.431448] device veth5110d03 entered promiscuous mode  
[142023.553919] eth0: renamed from vethad99f53  
[142023.566628] IPv6: ADDRCONF(NETDEV_CHANGE): veth5110d03: link becomes ready  
[142023.566665] docker0: port 1(veth5110d03) entered blocking state  
[142023.566667] docker0: port 1(veth5110d03) entered forwarding state  
[142346.809260] docker0: port 1(veth5110d03) entered disabled state  
[142346.812208] vethad99f53: renamed from eth0  
[142346.885705] docker0: port 1(veth5110d03) entered disabled state  
[142346.887918] device veth5110d03 left promiscuous mode  
[142346.887922] docker0: port 1(veth5110d03) entered disabled state  
[143930.554411] hello: loading out-of-tree module taints kernel.  
[143930.568461] hello: module verification failed: signature and/or required key  
missing - tainting kernel  
[143930.601743] Hello World!  
[143990.698431] Bye-bye World!.
```

Here when the module is inserted and removed, we monitor the kernel messages with sudo dmesg -k -w and observe the corresponding output.

Explanation:

In this task, we developed and tested a simple Loadable Kernel Module (LKM) to show how modules can be dynamically inserted into or removed from the Linux kernel without requiring a system reboot. The hello.c program outputs “Hello World!” when the module is loaded and “Bye-bye World!” when it is unloaded.

These messages are printed to the kernel log, which can be monitored using the dmesg command. To see this behavior, two terminals are used: one to continuously monitor kernel messages with sudo dmesg -k -w, and another to compile and manage the module using commands such as make, sudo insmod hello.ko (to insert the module), lsmod | grep hello (to verify it is loaded), and sudo rmmod hello (to remove it).

Task 1.B: Implement a Simple Firewall Using Netfilter

In this task, we will write our packet filtering program as an LKM, and then insert it into the packet processing path inside the kernel. This cannot be easily done in the past before netfilter was introduced into Linux.

Task:

1. Compile the code using the provided Makefile. Load it into the kernel, and demonstrate that the firewall is working as expected. You can use the following command to generate UDP packets to 8.8.8.8, which is Google's DNS server. If your firewall works, your request will be blocked; otherwise, you will get a response.

We can see that seedBlock is uncommented instead of seedFilter in the Makefile.

For the task we need seedFilter, hence we edit the Makefile

```
Makefile  seedBlock.c  seedFilter.c  seedPrint.c
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$cat Makefile
#obj-m += seedFilter.o
#obj-m += seedPrint.o
obj-m += seedBlock.o
all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules

clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean

ins:
    sudo dmesg -C
    sudo insmod seedFilter.ko

rm:
    sudo rmmod seedFilter

terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$nano Makefile
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$gedit Makefile
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
```

Step1: On the VM Terminal type the following command to make sure www.example.com is reachable, if it is not consider changing 8.8.8.8 to 8.8.4.4

- 3) Command: dig @8.8.8.8 www.example.com

VM:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
terminal1:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7
$dig @8.8.8.8 www.example.com

; <>> DiG 9.18.30-0ubuntu0.22.04.2-Ubuntu <>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 37827
;; flags: qr rd ra; QUERY: 1, ANSWER: 4, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 512
;; QUESTION SECTION:
;www.example.com.           IN      A

;; ANSWER SECTION:
www.example.com.      238      IN      CNAME   www.example.com-v4.edgesuite.net.
www.example.com-v4.edgesuite.net. 14567 IN CNAME a1422.dscr.akamai.net.
a1422.dscr.akamai.net. 20      IN      A       23.44.10.218
a1422.dscr.akamai.net. 20      IN      A       23.44.10.227

;; Query time: 43 msec
;; SERVER: 8.8.8.8#53(8.8.8.8) (UDP)
;; WHEN: Fri Oct 17 03:09:48 UTC 2025
;; MSG SIZE  rcvd: 154

terminal1:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7
$
```

This output shows a normal and successful dig command. A DNS query was sent to Google's server (8.8.8.8), which correctly responded with the IP addresses for www.example.com. This confirms that the network connection is functioning properly before the firewall is activated.

Step2: Open a new terminal to view kernel messages and execute the following -

- 4) Command: sudo dmesg -k -w

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
eration="profile_replace" profile="unconfined" name="snap.firefox.hook.install"
pid=119610 comm="apparmor_parser"
[149013.238130] audit: type=1400 audit(1760597885.711:116): apparmor="STATUS" op
eration="profile_replace" profile="unconfined" name="snap.firefox.hook.post-refr
esh" pid=119611 comm="apparmor_parser"
[150739.749866] e1000e 0000:02:00.0 ens160: NIC Link is Down
[151134.006497] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[157052.313815] e1000e 0000:02:00.0 ens160: NIC Link is Down
[157065.247208] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[168698.988008] e1000e 0000:02:00.0 ens160: NIC Link is Down
[168819.320644] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[168993.353389] e1000e 0000:02:00.0 ens160: NIC Link is Down
[169013.821960] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[169023.603086] e1000e 0000:02:00.0 ens160: NIC Link is Down
[169150.657967] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[169905.517159] e1000e 0000:02:00.0 ens160: NIC Link is Down
[169920.039333] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
```

- 5) Uncomment - 'obj-m += seedFilter.o' and comment the other two.

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$cat Makefile
obj-m += seedFilter.o
#obj-m += seedPrint.o
#obj-m += seedBlock.o
all:
        make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules

clean:
        make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean

ins:
        sudo dmesg -C
        sudo insmod seedFilter.ko

rm:
        sudo rmmod seedFilter

terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

- 6) Command:

```
$ make
```

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$make
make -C /lib/modules/5.15.0-153-generic/build M=/home/seed/Desktop/Lab7/volumes/
packet_filter modules
make[1]: Entering directory '/usr/src/linux-headers-5.15.0-153-generic'
warning: the compiler differs from the one used to build the kernel
  The kernel was built by: gcc (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0
  You are using:          gcc (Ubuntu 11.4.0-1ubuntu1~22.04.2) 11.4.0
  CC [M]  /home/seed/Desktop/Lab7/volumes/packet_filter/seedFilter.o
  MODPOST /home/seed/Desktop/Lab7/volumes/packet_filter/Module.symvers
  CC [M]  /home/seed/Desktop/Lab7/volumes/packet_filter/seedFilter.mod.o
  LD [M]  /home/seed/Desktop/Lab7/volumes/packet_filter/seedFilter.ko
  BTF [M] /home/seed/Desktop/Lab7/volumes/packet_filter/seedFilter.ko
Skipping BTF generation for /home/seed/Desktop/Lab7/volumes/packet_filter/seedFi
lter.ko due to unavailability of vmlinux
make[1]: Leaving directory '/usr/src/linux-headers-5.15.0-153-generic'
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

```
$ sudo insmod seedFilter.ko
```

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$ls
Makefile      seedBlock.c   seedFilter.mod   seedFilter.o
modules.order seedFilter.c   seedFilter.mod.c  seedPrint.c
Module.symvers seedFilter.ko  seedFilter.mod.o
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$sudo insmod seedFilter.ko
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

```
$ lsmod | grep seedFilter
```

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$lsmod | grep seedFilter
seedFilter           16384  0
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

terminal1:

```
[168993.353389] e1000e 0000:02:00.0 ens160: NIC Link is Down
[169013.821960] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[169023.603086] e1000e 0000:02:00.0 ens160: NIC Link is Down
[169150.657967] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[169905.517159] e1000e 0000:02:00.0 ens160: NIC Link is Down
[169920.039333] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[170524.828575] Registering filters.
[170548.037363] *** LOCAL_OUT
[170548.037560] 172.16.73.128 --> 34.107.243.93 (TCP)
[170548.044444] *** LOCAL_OUT
[170548.044464] 172.16.73.128 --> 34.107.243.93 (TCP)
[170616.312830] *** LOCAL_OUT
[170616.312847] 172.16.73.128 --> 8.8.8.8 (UDP)
[170616.312925] *** Dropping 8.8.8.8 (UDP), port 53
[170621.313636] *** LOCAL_OUT
[170621.313644] 172.16.73.128 --> 8.8.8.8 (UDP)
[170621.313716] *** Dropping 8.8.8.8 (UDP), port 53
[170626.321812] *** LOCAL_OUT
[170626.321865] 172.16.73.128 --> 8.8.8.8 (UDP)
[170626.322055] *** Dropping 8.8.8.8 (UDP), port 53
```

This dmesg log shows why the dig command timed out. The kernel is printing messages from your module, showing it's dropping the outgoing UDP packets destined for 8.8.8.8 on port 53 (the DNS port). The filter is correctly identifying and blocking the target traffic.

terminal2:

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$ dig @8.8.8.8 www.example.com
;; communications error to 8.8.8.8#53: timed out
;; communications error to 8.8.8.8#53: timed out
;; communications error to 8.8.8.8#53: timed out

; <>> DiG 9.18.30-0ubuntu0.22.04.2-Ubuntu <>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; no servers could be reached
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

Here, the seedFilter.ko module was successfully loaded using insmod and verified with lsmod. When the same dig command is executed afterward, it results in a “timed out” error, indicating that the firewall is now active and functioning as intended.

7) Remove the module by executing

Command: \$ sudo rmmod seedFilter

terminal1:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
[171271.276823] *** LOCAL_OUT
[171271.276830]      172.16.73.128  --> 151.101.1.91 (TCP)
[171271.292805] *** LOCAL_OUT
[171271.292812]      172.16.73.128  --> 151.101.1.91 (TCP)
[171271.298690] *** LOCAL_OUT
[171271.298694]      172.16.73.128  --> 34.107.243.93 (TCP)
[171323.282093] *** LOCAL_OUT
[171323.282184]      172.16.73.128  --> 34.107.243.93 (TCP)
[171323.290864] *** LOCAL_OUT
[171323.290872]      172.16.73.128  -> 34.107.243.93 (TCP)
[171323.291070] *** LOCAL_OUT
[171323.291073]      172.16.73.128  --> 34.107.243.93 (TCP)
[171323.295803] *** LOCAL_OUT
[171323.295807]      172.16.73.128  --> 34.107.243.93 (TCP)
[171324.288976] *** LOCAL_OUT
[171324.288989]      172.16.73.128  --> 151.101.1.91 (TCP)
[171324.290589] *** LOCAL_OUT
[171324.290598]      172.16.73.128  --> 151.101.1.91 (TCP)
[171324.290784] *** LOCAL_OUT
[171324.290788]      172.16.73.128  --> 151.101.1.91 (TCP)
[171324.294645] *** LOCAL_OUT
[171324.294650]      172.16.73.128  --> 151.101.1.91 (TCP)
[171415.656021] The filters are being removed.
```

terminal2:

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$sudo rmmod seedFilter
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$lsmod | grep seedFilter
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

8) To clear the kernel messages execute -

Command: \$ dmesg -C

```
terminal1:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7
$sudo dmesg -C
terminal1:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7
$sudo dmesg -k -w
```

Step3: Lets retry running dig

9) Command: dig @8.8.8.8 www.example.com

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
terminal1:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7
$dig @8.8.8.8 www.example.com

; <>> DiG 9.18.30-0ubuntu0.22.04.2-Ubuntu <>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 37827
;; flags: qr rd ra; QUERY: 1, ANSWER: 4, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 512
;; QUESTION SECTION:
;www.example.com.           IN      A

;; ANSWER SECTION:
www.example.com.      238     IN      CNAME   www.example.com-v4.edgesuite.net.
www.example.com-v4.edgesuite.net. 14567 IN CNAME a1422.dscr.akamai.net.
a1422.dscr.akamai.net. 20      IN      A       23.44.10.218
a1422.dscr.akamai.net. 20      IN      A       23.44.10.227

;; Query time: 43 msec
;; SERVER: 8.8.8.8#53(8.8.8.8) (UDP)
;; WHEN: Fri Oct 17 03:09:48 UTC 2025
;; MSG SIZE  rcvd: 154

terminal1:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7
$
```

Running the dig command works normally, as expected

Explanation:

In this task, we implemented a simple firewall using the Linux Netfilter framework by loading a custom kernel module (LKM). The module, `seedFilter.ko`, filters outgoing UDP packets destined for Google's DNS server (8.8.8.8) on port 53. Before enabling the firewall, a DNS lookup using the dig command works normally, confirming internet connectivity. After editing the Makefile to enable the firewall module and successfully compiling and inserting it into the kernel, the firewall begins intercepting and dropping the targeted packets. This behavior is confirmed through kernel logs in `dmesg`, where the dropped packets are reported, and through the dig command timing out, showing that DNS requests are blocked as intended. Finally, removing the module restores normal functionality, allowing DNS packets to pass again.

Task:

2. Hook the printInfo function to all of the netfilter hooks. Here are the macros of the hook numbers. Using your experiment results to help explain at what condition each of the hook functions be invoked.
 - a. NF_INET_PRE_ROUTING
 - b. NF_INET_LOCAL_IN
 - c. NF_INET_FORW ARD
 - d. NF_INET_LOCAL_OUT
 - e. NF_INET_POST_ROUTING

Step1: Change the code for this task

- 10) Uncomment - 'obj-m += seedPrint.o' and comment the other two.

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$gedit Makefile
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$cat Makefile
#obj-m += seedFilter.o
obj-m += seedPrint.o
#obj-m += seedBlock.o
all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules

clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean

ins:
    sudo dmesg -C
    sudo insmod seedFilter.ko

rm:
    sudo rmmod seedFilter

terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

Step2: Open a new terminal to view kernel messages and execute the following -

- 11) Command: sudo dmesg -k -w

```
terminal1:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7
$sudo dmesg -k -w
[171660.399587] e1000e 0000:02:00.0 ens160: NIC Link is Down
[171674.242002] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[172161.913594] device br-b277b226ad85 entered promiscuous mode
[172192.580423] device br-b277b226ad85 left promiscuous mode
[172200.360087] device br-b277b226ad85 entered promiscuous mode
[172281.318177] device br-b277b226ad85 left promiscuous mode
[172878.860747] device br-b277b226ad85 entered promiscuous mode
```

(Here i opened wireshark just to monitor hence the promiscuous mode)

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

12)Command:

```
$ make
```

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$make
make -C /lib/modules/5.15.0-153-generic/build M=/home/seed/Desktop/Lab7/volumes/
packet_filter modules
make[1]: Entering directory '/usr/src/linux-headers-5.15.0-153-generic'
warning: the compiler differs from the one used to build the kernel
  The kernel was built by: gcc (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0
  You are using:           gcc (Ubuntu 11.4.0-1ubuntu1~22.04.2) 11.4.0
  CC [M]  /home/seed/Desktop/Lab7/volumes/packet_filter/seedPrint.o
  MODPOST /home/seed/Desktop/Lab7/volumes/packet_filter/Module.symvers
  CC [M]  /home/seed/Desktop/Lab7/volumes/packet_filter/seedPrint.mod.o
  LD [M]  /home/seed/Desktop/Lab7/volumes/packet_filter/seedPrint.ko
  BTF [M] /home/seed/Desktop/Lab7/volumes/packet_filter/seedPrint.ko
Skipping BTF generation for /home/seed/Desktop/Lab7/volumes/packet_filter/seedPr
int.ko due to unavailability of vmlinux
make[1]: Leaving directory '/usr/src/linux-headers-5.15.0-153-generic'
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

```
$ sudo insmod seedPrint.ko
```

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$ls
Makefile      seedFilter.c      seedFilter.mod.o  seedPrint.mod
modules.order  seedFilter.ko     seedFilter.o      seedPrint.mod.c
Module.symvers seedFilter.mod    seedPrint.c      seedPrint.mod.o
seedBlock.c    seedFilter.mod.c  seedPrint.ko     seedPrint.o
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$sudo insmod seedPrint.ko
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
```

```
$ lsmod | grep seedPrint
```

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$lsmod | grep seedPrint
seedPrint          16384  0
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

terminal1:

```
terminal1:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7
$sudo dmesg -k -w
[171660.399587] e1000e 0000:02:00.0 ens160: NIC Link is Down
[171674.242002] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[172161.913594] device br-b277b226ad85 entered promiscuous mode
[172192.580423] device br-b277b226ad85 left promiscuous mode
[172200.360087] device br-b277b226ad85 entered promiscuous mode
[172281.318177] device br-b277b226ad85 left promiscuous mode
[172878.860747] device br-b277b226ad85 entered promiscuous mode
[173192.383821] Registering filters.
```

13)Now run the below command

Command:

```
dig @8.8.8.8 www.example.com
```

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

terminal2:

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$dig @8.8.8.8 www.example.com

; <>> DiG 9.18.30-0ubuntu0.22.04.2-Ubuntu <>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 44346
;; flags: qr rd ra; QUERY: 1, ANSWER: 4, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags:; udp: 512
;; QUESTION SECTION:
;www.example.com.           IN      A

;; ANSWER SECTION:
www.example.com.      27      IN      CNAME   www.example.com-v4.edgesuite.net.
www.example.com-v4.edgesuite.net. 6220 IN CNAME a1422.dscr.akamai.net.
a1422.dscr.akamai.net. 20      IN      A       23.44.10.227
a1422.dscr.akamai.net. 20      IN      A       23.44.10.218

;; Query time: 48 msec
;; SERVER: 8.8.8.8#53(8.8.8.8) (UDP)
;; WHEN: Fri Oct 17 04:08:43 UTC 2025
;; MSG SIZE rcvd: 154

terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

This time, running the dig command works normally, as expected, since this module is designed to print output rather than block execution.

terminal1:

```
[173313.089056]      127.0.0.1  --> 224.0.0.251 (UDP)
[173313.089116] *** PRE_ROUTING
[173313.089117]      127.0.0.1  --> 224.0.0.251 (UDP)
[173313.089123] *** LOCAL_IN
[173313.089125]      127.0.0.1  --> 224.0.0.251 (UDP)
[173313.089801] *** LOCAL_OUT
[173313.089804]      10.9.0.1  --> 224.0.0.251 (UDP)
[173313.089828] *** POST_ROUTING
[173313.089829]      10.9.0.1  --> 224.0.0.251 (UDP)
[173313.089870] *** PRE_ROUTING
[173313.089871]      10.9.0.1  --> 224.0.0.251 (UDP)
[173313.089875] *** LOCAL_IN
[173313.089875]      10.9.0.1  --> 224.0.0.251 (UDP)
[173313.089889] *** POST_ROUTING
[173313.089891]      10.9.0.1  --> 224.0.0.251 (UDP)
[173342.467119] *** LOCAL_OUT
[173342.467136]      172.16.73.128 --> 8.8.8.8 (UDP)
[173342.467266] *** POST_ROUTING
[173342.467266]      172.16.73.128 --> 8.8.8.8 (UDP)
[173342.515902] *** PRE_ROUTING
[173342.515906]      8.8.8.8 --> 172.16.73.128 (UDP)
[173342.515936] *** LOCAL_IN
[173342.515937]      8.8.8.8 --> 172.16.73.128 (UDP)
```

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

14) Remove the module by executing

Command: \$ sudo rmmod seedPrint

terminal1:

```
[173313.089116] *** PRE_ROUTING
[173313.089117]      127.0.0.1 --> 224.0.0.251 (UDP)
[173313.089123] *** LOCAL_IN
[173313.089125]      127.0.0.1 --> 224.0.0.251 (UDP)
[173313.089801] *** LOCAL_OUT
[173313.089804]      10.9.0.1 --> 224.0.0.251 (UDP)
[173313.089828] *** POST_ROUTING
[173313.089829]      10.9.0.1 --> 224.0.0.251 (UDP)
[173313.089870] *** PRE_ROUTING
[173313.089871]      10.9.0.1 --> 224.0.0.251 (UDP)
[173313.089875] *** LOCAL_IN
[173313.089875]      10.9.0.1 --> 224.0.0.251 (UDP)
[173313.089889] *** POST_ROUTING
[173313.089891]      10.9.0.1 --> 224.0.0.251 (UDP)
[173342.467119] *** LOCAL_OUT
[173342.467136]      172.16.73.128 --> 8.8.8.8 (UDP)
[173342.467266] *** POST_ROUTING
[173342.467266]      172.16.73.128 --> 8.8.8.8 (UDP)
[173342.515902] *** PRE_ROUTING
[173342.515906]      8.8.8.8 --> 172.16.73.128 (UDP)
[173342.515936] *** LOCAL_IN
[173342.515937]      8.8.8.8 --> 172.16.73.128 (UDP)
[173445.593726] The filters are being removed.
```

terminal2:

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$ sudo rmmod seedPrint
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$ lsmod | grep seedPrint
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

Wireshark: (no useful packet capture)

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

No.	Time	Source	Destination	Protocol	Length	Info
1	2025-10-17 04:0...	10.9.0.1	224.0.0.251	MDNS	104	Standard query 0x0000 PTR _ipps._tcp.local, "QM" question PTR _nmea-0183._tcp.local
2	2025-10-17 04:0...	fe80::e063:17ff:fe4...	ff02::fb	MDNS	124	Standard query 0x0000 PTR _ipps._tcp.local, "QM" question PTR _nmea-0183._tcp.local

Explanation:

In this task, we used the Netfilter framework to observe how packets move through different stages of the Linux networking stack. The seedPrint.ko kernel module was modified to attach the printInfo function to all five Netfilter hook points:

- NF_INET_PRE_ROUTING: Packets arriving from the network before routing decisions
- NF_INET_LOCAL_IN: Packets destined for the local machine
- NF_INET_FORWARD: Packets being routed through the machine to another destination
- NF_INET_LOCAL_OUT: Outgoing packets generated by the local system
- NF_INET_POST_ROUTING: Packets leaving the system after routing is completed

After loading the module, the dig command was run again. Since this module does not filter or drop traffic, DNS requests proceeded normally. However, kernel messages showed logging from each hook, demonstrating exactly where packets are intercepted during transmission and reception. Removing the module stopped the logging and returned the system to normal.

This experiment helps visualize the Linux packet processing path and understand when each Netfilter hook gets invoked based on the packet's direction and destination.

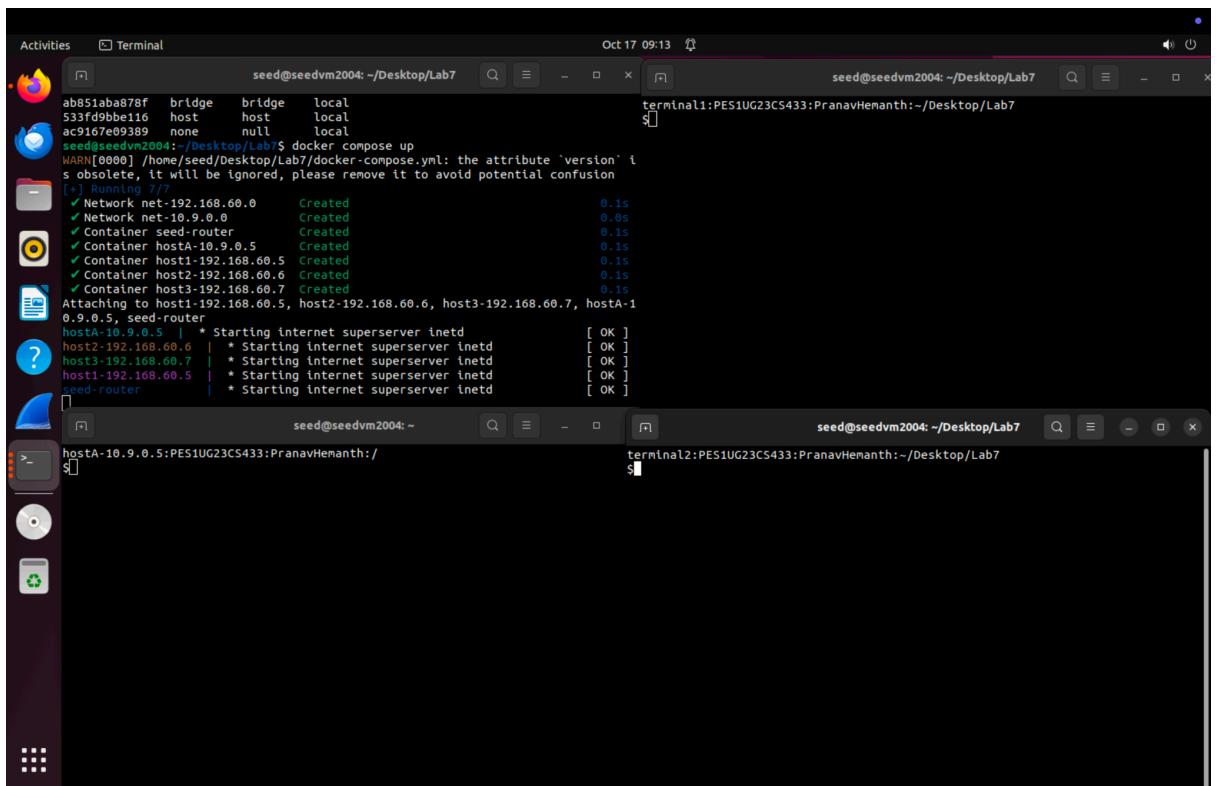
Task:

3. Implement two more hooks to achieve the following:

- a. preventing other computers to ping the VM, and
- b. preventing other computers from telnetting into the VM.

Please implement two different hook functions, but register them to the same netfilter hook. You should decide what hook to use. Telnet's default port is TCP port 23. To test it, you can start the containers, go to 10.9.0.5, run the following commands (10.9.0.1 is the IP address assigned to the VM; for the sake of simplicity, you can hardcode this IP address in your firewall rules

Three window setup:



Step1: Open a new terminal to view kernel messages and execute the following -

15)Command: sudo dmesg -k -w

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
[180952.882402] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[181008.583444] e1000e 0000:02:00.0 ens160: NIC Link is Down
[181264.641181] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[182336.518425] e1000e 0000:02:00.0 ens160: NIC Link is Down
[182361.477063] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[182628.943116] e1000e 0000:02:00.0 ens160: NIC Link is Down
[182636.052154] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[182730.232006] e1000e 0000:02:00.0 ens160: NIC Link is Down
[182743.840793] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[183016.578784] e1000e 0000:02:00.0 ens160: NIC Link is Down
[183064.312569] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[183104.645321] e1000e 0000:02:00.0 ens160: NIC Link is Down
[183117.128770] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[183122.241320] e1000e 0000:02:00.0 ens160: NIC Link is Down
[183127.369041] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
```

16)

Step2: Change the code for this task

17) Uncomment - 'obj-m += seedBlock.o' and comment the other two.

```
1 #obj-m += seedFilter.o
2 #obj-m += seedPrint.o
3 obj-m += seedBlock.o
4 all:
5     make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules
6
7 clean:
8     make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean
9
10 ins:
11     sudo dmesg -C
12     sudo insmod seedFilter.ko
13
14 rm:
15     sudo rmmod seedFilter
16
```

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$gedit Makefile
^C
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$cat Makefile
#obj-m += seedFilter.o
#obj-m += seedPrint.o
obj-m += seedBlock.o
all:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) modules

clean:
    make -C /lib/modules/$(shell uname -r)/build M=$(PWD) clean

ins:
    sudo dmesg -c
    sudo insmod seedFilter.ko

rm:
    sudo rmmod seedFilter

terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

18)Command:

```
$ make
```

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$make
make -C /lib/modules/5.15.0-153-generic/build M=/home/seed/Desktop/Lab7/volumes/
packet_filter modules
make[1]: Entering directory '/usr/src/linux-headers-5.15.0-153-generic'
warning: the compiler differs from the one used to build the kernel
  The kernel was built by: gcc (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0
  You are using:          gcc (Ubuntu 11.4.0-1ubuntu1~22.04.2) 11.4.0
  CC [M]  /home/seed/Desktop/Lab7/volumes/packet_filter/seedBlock.o
  MODPOST /home/seed/Desktop/Lab7/volumes/packet_filter/Module.symvers
  CC [M]  /home/seed/Desktop/Lab7/volumes/packet_filter/seedBlock.mod.o
  LD [M]  /home/seed/Desktop/Lab7/volumes/packet_filter/seedBlock.ko
  BTF [M] /home/seed/Desktop/Lab7/volumes/packet_filter/seedBlock.ko
Skipping BTF generation for /home/seed/Desktop/Lab7/volumes/packet_filter/seedBl
ock.ko due to unavailability of vmlinux
make[1]: Leaving directory '/usr/src/linux-headers-5.15.0-153-generic'
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

```
$ sudo insmod seedPrint.ko
```

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$ls
Makefile      seedBlock.mod   seedFilter.ko     seedPrint.c      seedPrint.o
modules.order  seedBlock.mod.c  seedFilter.mod   seedPrint.ko
Module.symvers seedBlock.mod.o  seedFilter.mod.c  seedPrint.mod
seedBlock.c    seedBlock.o     seedFilter.mod.o  seedPrint.mod.c
seedBlock.ko    seedFilter.c    seedFilter.o     seedPrint.mod.o
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$sudo insmod seedBlock.ko
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
$ lsmod | grep seedPrint
```

```
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$lsmod | grep seedBlock
seedBlock           16384  0
terminal2:PES1UG23CS433:PranavHemanth:~/Desktop/Lab7/volumes/packet_filter
$
```

terminal1:

```
[182336.518425] e1000e 0000:02:00.0 ens160: NIC Link is Down
[182361.477063] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[182628.943116] e1000e 0000:02:00.0 ens160: NIC Link is Down
[182636.052154] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[182730.232006] e1000e 0000:02:00.0 ens160: NIC Link is Down
[182743.840793] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[183016.578784] e1000e 0000:02:00.0 ens160: NIC Link is Down
[183064.312569] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[183104.645321] e1000e 0000:02:00.0 ens160: NIC Link is Down
[183117.128770] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[183122.241320] e1000e 0000:02:00.0 ens160: NIC Link is Down
[183127.369041] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[183406.197561] Registering filters.          []
[183429.679860] *** LOCAL_OUT
[183429.680071]      172.16.73.128  --> 34.107.243.93 (TCP)
[183429.690897] *** LOCAL_OUT
[183429.690904]      172.16.73.128  --> 34.107.243.93 (TCP)
```

19)Now run the below command on Host A - 10.9.0.5

Command:

```
$ ping 10.9.0.1
```

hostA-10.9.0.5:

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/
$ping 10.9.0.1
PING 10.9.0.1 (10.9.0.1) 56(84) bytes of data.
^C
--- 10.9.0.1 ping statistics ---
9 packets transmitted, 0 received, 100% packet loss, time 8176ms

hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/
$
```

terminal1:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
, Flow Control: None
[183104.645321] e1000e 0000:02:00.0 ens160: NIC Link is Down
[183117.128770] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[183122.241320] e1000e 0000:02:00.0 ens160: NIC Link is Down
[183127.369041] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex
, Flow Control: None
[183406.197561] Registering filters.
[183429.679860] *** LOCAL_OUT
[183429.680071]      172.16.73.128  --> 34.107.243.93 (TCP)
[183429.690897] *** LOCAL_OUT
[183429.690904]      172.16.73.128  --> 34.107.243.93 (TCP)
[183481.828593] *** LOCAL_OUT
[183481.828666]      172.16.73.128  --> 185.125.190.57 (UDP)
[183548.868621] *** Dropping 10.9.0.1 (ICMP)
[183549.873894] *** Dropping 10.9.0.1 (ICMP)
[183550.897887] *** Dropping 10.9.0.1 (ICMP)
[183551.922409] *** Dropping 10.9.0.1 (ICMP)
[183552.949627] *** Dropping 10.9.0.1 (ICMP)
[183553.969897] *** Dropping 10.9.0.1 (ICMP)           ]
[183554.993591] *** Dropping 10.9.0.1 (ICMP)
[183556.017778] *** Dropping 10.9.0.1 (ICMP)
[183557.042802] *** Dropping 10.9.0.1 (ICMP)
]
```

\$ telnet 10.9.0.1

hostA-10.9.0.5:

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/ 
$telnet 10.9.0.1
Trying 10.9.0.1...
^C
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/ 
$
```

terminal1:

```
[183127.369041] e1000e 0000:02:00.0 ens160: NIC Link is Up 1000 Mbps Full Duplex, Flow Control: None
[183406.197561] Registering filters.
[183429.679860] *** LOCAL_OUT
[183429.680071]     172.16.73.128 --> 34.107.243.93 (TCP)
[183429.690897] *** LOCAL_OUT
[183429.690904]     172.16.73.128 --> 34.107.243.93 (TCP)
[183481.828593] *** LOCAL_OUT
[183481.828666]     172.16.73.128 --> 185.125.190.57 (UDP)
[183548.868621] *** Dropping 10.9.0.1 (ICMP)
[183549.873894] *** Dropping 10.9.0.1 (ICMP)
[183550.897887] *** Dropping 10.9.0.1 (ICMP)
[183551.922409] *** Dropping 10.9.0.1 (ICMP)
[183552.949627] *** Dropping 10.9.0.1 (ICMP)
[183553.969897] *** Dropping 10.9.0.1 (ICMP)
[183554.993591] *** Dropping 10.9.0.1 (ICMP)
[183556.017778] *** Dropping 10.9.0.1 (ICMP)    []
[183557.042802] *** Dropping 10.9.0.1 (ICMP)
[183638.693627] *** LOCAL_OUT
[183638.693731]     172.16.73.128 --> 224.0.0.251 (UDP)
[183673.006888] *** Dropping 10.9.0.1 (TCP), port 23
[183674.023231] *** Dropping 10.9.0.1 (TCP), port 23
[183676.038793] *** Dropping 10.9.0.1 (TCP), port 23
[183680.230520] *** Dropping 10.9.0.1 (TCP), port 23
```

Explanation:

In this task, we extended the Netfilter-based firewall to block two specific types of incoming traffic: ICMP echo requests (ping) and TCP connections to port 23 (telnet). We created two separate filtering functions but registered both to the same Netfilter hook, since both packet types are inbound and should be dropped before reaching the system. After compiling and loading the seedBlock.ko module, the VM became protected from external ping and telnet requests.

Testing from another host (10.9.0.5) confirmed the behavior. Ping attempts to the VM's IP address (10.9.0.1) resulted in 100 percent packet loss, and telnet attempts hung without establishing a connection. Meanwhile, kernel log messages showed that the firewall was detecting and dropping ICMP packets and TCP packets destined for port 23. This validates that the firewall rules are correctly identifying and blocking the targeted traffic.

The experiment demonstrates how custom Netfilter hooks can enforce basic firewall policies and prevent unauthorized access to critical services on the VM.

Task 2: Experimenting with Stateless Firewall Rules

In the previous task, we had a chance to build a simple firewall using netfilter. Actually, Linux already has a built-in firewall, also based on netfilter. This firewall is called iptables. Technically, the kernel part implementation of the firewall is called Xtables, while iptables is a user-space program to configure the firewall. However, iptables is often used to refer to both the kernel-part implementation and the user-space program.

In this task, we will use iptables to set up a firewall. The iptables firewall is designed not only to filter packets, but also to make changes to packets. To help manage these firewall rules for different purposes, iptables organizes all rules using a hierarchical structure: table, chain, and rules. There are several tables, each specifying the main purpose of the rules as shown in Table 1. For example, rules for packet filtering should be placed in the filter table, while rules for making changes to packets should be placed in the nat or mangle tables.

Task 2.A: Protecting the Router

Step1: Ensure all containers are running

```
seed@seedvm2004:~$ docker ps
CONTAINER ID   IMAGE           COMMAND
CREATED        STATUS          PORTS      NAMES
dcca3a58e38   seed-router-image   "bash -c ' ip route ..."
6 days ago    Up 4 minutes   seed-router
344d644c5360  handsonsecurity/seed-ubuntu:large-arm   "bash -c ' ip route ..."
6 days ago    Up 4 minutes   host3-192.168.60.7
44e8f7e115a1  handsonsecurity/seed-ubuntu:large-arm   "bash -c ' ip route ..."
6 days ago    Up 4 minutes   host1-192.168.60.5
64850cb17f44  handsonsecurity/seed-ubuntu:large-arm   "bash -c ' ip route ..."
6 days ago    Up 4 minutes   hostA-10.9.0.5
a57272dc8a30  handsonsecurity/seed-ubuntu:large-arm   "bash -c ' ip route ..."
6 days ago    Up 4 minutes   host2-192.168.60.6
seed@seedvm2004:~$
```

Step2: In order to view the current policies run the below command (on seed router)
 20)Command: iptables -t filter -L -n

```
seed-router:PES1UG23CS433:PranavHemanth:/
$ iptables -t filter -L -n
Chain INPUT (policy ACCEPT)
target     prot opt source          destination
target     prot opt source          destination

Chain FORWARD (policy ACCEPT)
target     prot opt source          destination
target     prot opt source          destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source          destination
seed-router:PES1UG23CS433:PranavHemanth:/
$
```

Above are initial configs

Step3: Now run the below on seed router and try to access from hostA

21)Command:

```
iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEPT
iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT
iptables -P OUTPUT DROP
iptables -P INPUT DROP
iptables -t filter -L -n
```

seed-router:

```
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEPT
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -P OUTPUT DROP
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -P INPUT DROP
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -t filter -L -n
Chain INPUT (policy DROP)
target      prot opt source          destination
ACCEPT     icmp  --  0.0.0.0/0        0.0.0.0/0          icmp type 8
ACCEPT     icmp  --  0.0.0.0/0        0.0.0.0/0          icmp type 8

Chain FORWARD (policy ACCEPT)
target      prot opt source          destination

Chain OUTPUT (policy DROP)
target      prot opt source          destination
ACCEPT     icmp  --  0.0.0.0/0        0.0.0.0/0          icmp type 0
seed-router:PES1UG23CS433:PranavHemanth:/
```

Step4: Now try to access (ping and telnet) the router from Host A - 10.9.0.5

22)Command:

```
ping seed-router
```

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$ ping seed-router  
PING seed-router (10.9.0.11) 56(84) bytes of data.  
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=1 ttl=64 time=0.40  
ms  
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=2 ttl=64 time=0.16  
ms  
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=3 ttl=64 time=0.14  
ms  
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=4 ttl=64 time=0.14  
ms  
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=5 ttl=64 time=0.13  
ms  
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=6 ttl=64 time=0.11  
ms  
^C  
--- seed-router ping statistics ---  
6 packets transmitted, 6 received, 0% packet loss, time 5102ms  
rtt min/avg/max/mdev = 0.115/0.184/0.401/0.098 ms  
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$
```

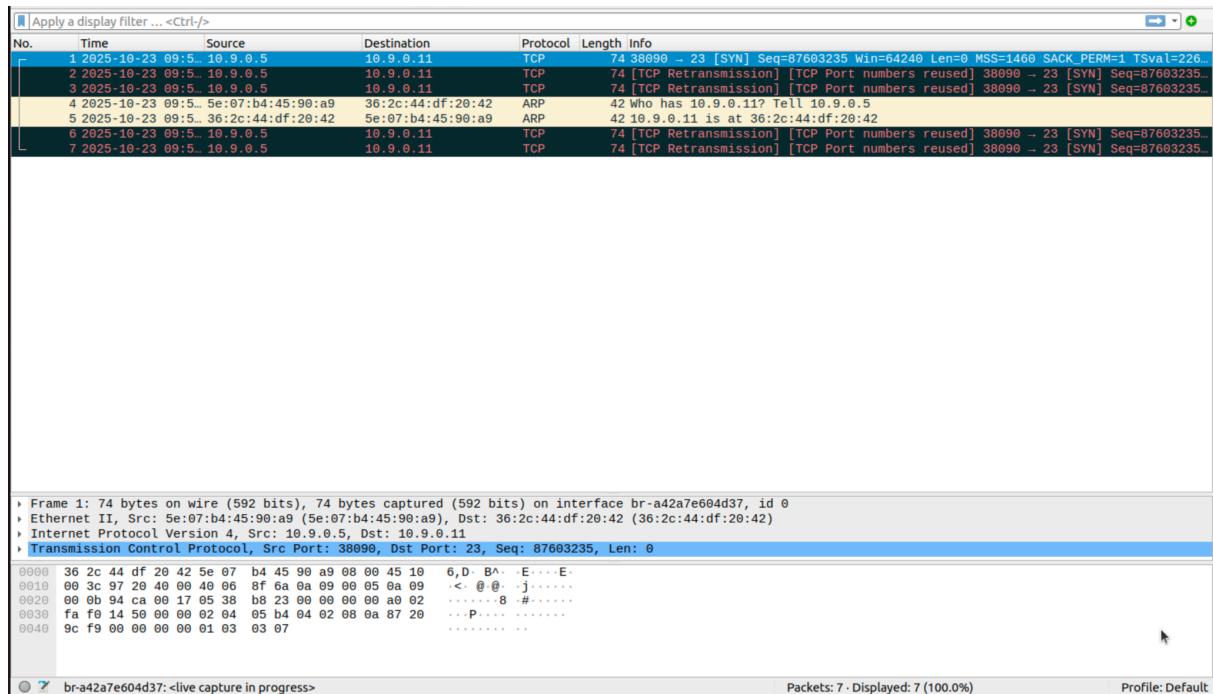
Wireshark:

No.	Time	Source	Destination	Protocol	Length	Info
1	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=1/256, ttl=64 (reply in 2)
2	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=1/256, ttl=64 (request in 1)
3	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=2/512, ttl=64 (reply in 4)
4	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=2/512, ttl=64 (request in 3)
5	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=3/768, ttl=64 (reply in 6)
6	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=3/768, ttl=64 (request in 5)
7	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=4/1024, ttl=64 (reply in 8)
8	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=4/1024, ttl=64 (request in 7)
9	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=5/1280, ttl=64 (reply in 10)
10	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=5/1280, ttl=64 (request in 9)
11	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=6/1536, ttl=64 (reply in 12)
12	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=6/1536, ttl=64 (request in 11)
13	2025-10-23 09:5. 36:2c:44:df:20:42	5e:07:b4:45:90:a9		ARP	42	Who has 10.9.0.5? Tell 10.9.0.11
14	2025-10-23 09:5. 5e:07:b4:45:90:a9	36:2c:44:df:20:42		ARP	42	Who has 10.9.0.11? Tell 10.9.0.5
15	2025-10-23 09:5. 5e:07:b4:45:90:a9	36:2c:44:df:20:42		ARP	42	10.9.0.5 is at 36:2c:44:df:20:42
16	2025-10-23 09:5. 36:2c:44:df:20:42	5e:07:b4:45:90:a9		ARP	42	10.9.0.11 is at 36:2c:44:df:20:42
17	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=7/1792, ttl=64 (reply in 18)
18	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=7/1792, ttl=64 (request in 17)
19	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=8/2048, ttl=64 (reply in 20)
20	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=8/2048, ttl=64 (request in 19)
21	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) request id=0x0002, seq=9/2304, ttl=64 (reply in 22)
22	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=9/2304, ttl=64 (request in 21)
23	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=10/2560, ttl=64 (reply in 24)
24	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=10/2560, ttl=64 (request in 23)
25	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=11/2816, ttl=64 (reply in 26)
26	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=11/2816, ttl=64 (request in 25)
27	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=12/3072, ttl=64 (reply in 28)
28	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=12/3072, ttl=64 (request in 27)
29	2025-10-23 09:5. 10. 9. 0.5	10. 9. 0.11		ICMP	98	Echo (ping) request id=0x0002, seq=13/3328, ttl=64 (reply in 30)
30	2025-10-23 09:5. 10. 9. 0.11	10. 9. 0.5		ICMP	98	Echo (ping) reply id=0x0002, seq=13/3328, ttl=64 (request in 30)
Frame: 1: 98 bytes on wire (784 bits), 98 bytes captured (784 bits) on interface br-a42a7e604d37, id 0						
> Ethernet II, Src: 5e:07:b4:45:90:a9 (5e:07:b4:45:90:a9), Dst: 36:2c:44:df:20:42 (36:2c:44:df:20:42)						
Internet Protocol Version 4, Src: 10.9.0.5, Dst: 10.9.0.11						
> Internet Control Message Protocol						
0000	36 2c 44 df 20 42 5e 07	b4 45 90 a9 08 00 45 00	6, D B A E E E E			
0010	00 54 8d e1 40 00 01 91	98 a0 09 00 05 00 09	-T @			
0020	00 0b 08 77 82 00 02	00 01 fa f9 08 00 00	... w h			
0030	00 00 3e 44 08 00 00 00	00 00 10 11 12 13 14 15	>D			
0040	16 17 18 19 1a 1b 1c 1d	1e 1f 20 21 22 23 24 25 !%#%			
0050	26 27 28 29 2a 2b 2c 2d	2e 2f 30 31 32 33 34 35	'(' +, . /012345			
0060	36 37		67			

telnet seed-router

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$telnet seed-router  
Trying 10.9.0.11...  
^C  
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$
```

Wireshark:



Questions :

(1) Can you ping the router?

Answer: Yes. The ping succeeds with 0 percent packet loss, because ICMP echo-request and echo-reply packets were explicitly allowed in the firewall rules.

(2) Can you telnet into the router (a telnet server is running on all the containers; an account called seed was created on them with a password dees).

Answer: No. The telnet connection fails and eventually times out, because TCP packets (port 23) are not allowed by any rule and are therefore dropped by the default INPUT and OUTPUT DROP policies.

Step5: Cleanup

23)Command:

```
iptables -F
iptables -P OUTPUT ACCEPT
iptables -P INPUT ACCEPT
```

```
seed-router:PES1UG23CS433:PranavHemanth:/  
$ iptables -F  
seed-router:PES1UG23CS433:PranavHemanth:/  
$ iptables -P OUTPUT ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$ iptables -P INPUT ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$ iptables -t filter -L -n  
Chain INPUT (policy ACCEPT)  
target     prot opt source          destination  
Chain FORWARD (policy ACCEPT)  
target     prot opt source          destination  
Chain OUTPUT (policy ACCEPT)  
target     prot opt source          destination  
seed-router:PES1UG23CS433:PranavHemanth:/  
$ █
```

Another way to restore the states of all the tables is to restart the container.

You can do it using the following command (you need to find the container's ID first):

```
$ docker restart <container ID> (Here - $ docker restart dccca3a58e38)
```

Explanation:

In this task, we explored Linux's built-in firewall, iptables, to protect the router by enforcing stateless filtering rules. Initially, all firewall chains (INPUT, OUTPUT, FORWARD) had default ACCEPT policies, which allowed unrestricted traffic. We then modified the router's firewall behavior to adopt a default-deny stance by setting both INPUT and OUTPUT chains to DROP, while adding specific exceptions to allow ICMP ping traffic:

- Allow incoming ICMP echo-requests
- Allow outgoing ICMP echo-replies

After applying these rules, we tested connectivity from Host A. Ping requests to the seed-router were successful because the inserted ACCEPT rules specifically permitted ICMP request and reply packets. However, telnet attempts failed because TCP packets do not match any ACCEPT rule and were instead dropped by the default DROP policies. Packet captures confirmed that telnet traffic was silently discarded, while ping traffic continued to pass through correctly.

This experiment demonstrates how stateless firewall rules operate, how rule order and default policies influence traffic decisions, and how selective protocol-based restrictions can secure a networked device against unauthorized access.

Task 2.B: Protecting the Internal Network

In this task, we will set up firewall rules on the router to protect the internal network 192.168.60.0/24. We need to use the FORWARD chain for this purpose.

In this task, we want to implement a firewall to protect the internal network. More specifically, we need to enforce the following restrictions on the ICMP traffic:

1. Outside hosts cannot ping internal hosts.
2. Outside hosts can ping the router.
3. Internal hosts can ping outside hosts.
4. All other packets between the internal and external networks should be blocked.

Step1: Execute the following iptables commands on the seed-router container-24)Command:

```
iptables -A FORWARD -i eth0 -p icmp --icmp-type echo-request -j DROP
iptables -A FORWARD -i eth1 -p icmp --icmp-type echo-request -j ACCEPT
iptables -A FORWARD -i eth0 -p icmp --icmp-type echo-reply -j ACCEPT
iptables -P FORWARD DROP
iptables -L -n -v
```

seed-router:

```
seed-router:PES1UG23CS433:PranavHemanth:/
$ iptables -A FORWARD -i eth0 -p icmp --icmp-type echo-request -j DROP
seed-router:PES1UG23CS433:PranavHemanth:/
$ iptables -A FORWARD -i eth1 -p icmp --icmp-type echo-request -j ACCEPT
seed-router:PES1UG23CS433:PranavHemanth:/
$ iptables -A FORWARD -i eth0 -p icmp --icmp-type echo-reply -j ACCEPT
seed-router:PES1UG23CS433:PranavHemanth:/
$ iptables -P FORWARD DROP
seed-router:PES1UG23CS433:PranavHemanth:/
$ iptables -L -n -v
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source               destination
Chain FORWARD (policy DROP 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source               destination
  0     0   DROP      icmp  --  eth0    *      0.0.0.0/0            0.0.0.0/0          icmp type 8
  0     0   DROP      icmp  --  eth0    *      0.0.0.0/0            0.0.0.0/0          icmp type 8
  0     0   DROP      icmp  --  eth0    *      0.0.0.0.0/0         0.0.0.0.0/0        icmp type 8
  0     0   ACCEPT    icmp  --  eth1    *      0.0.0.0/0            0.0.0.0/0          icmp type 8
  0     0   ACCEPT    icmp  --  eth0    *      0.0.0.0/0            0.0.0.0/0          icmp type 0
Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source               destination
seed-router:PES1UG23CS433:PranavHemanth:/
$
```

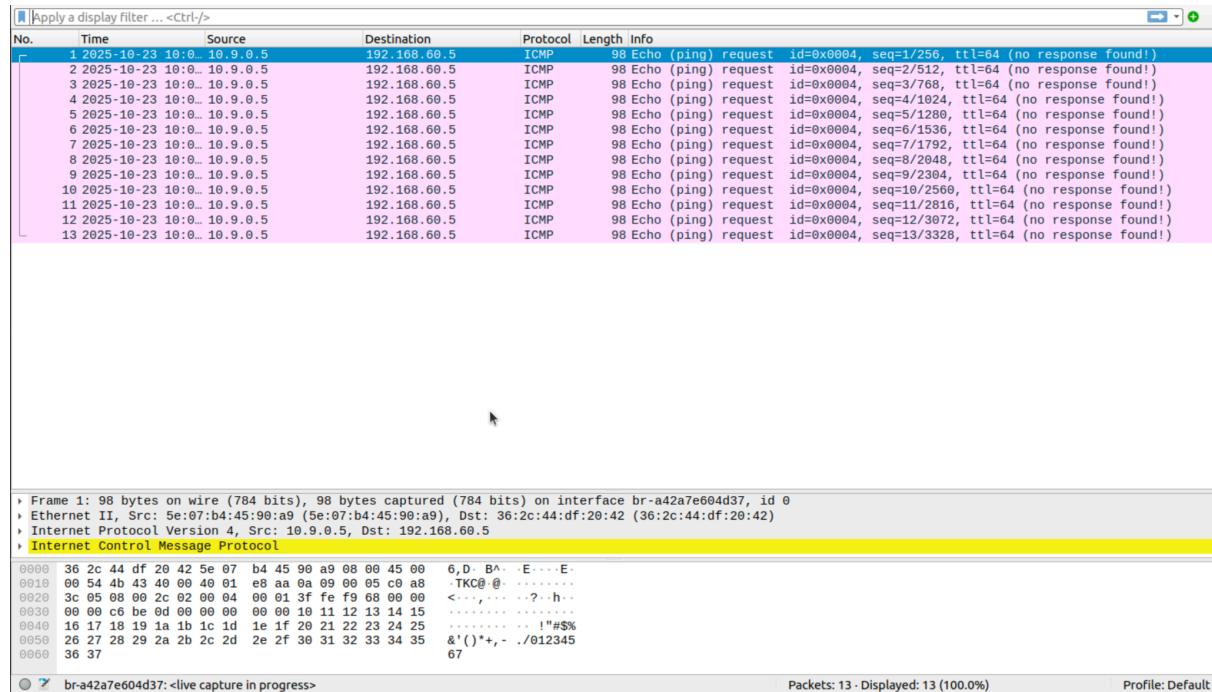
Step2: Now we shall see if these restrictions have been enforced in the network.

1. Outside hosts cannot ping internal hosts.
 - a. Command: ping 192.168.60.5 (On Host A - 10.9.0.5)

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$ping 192.168.60.5  
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.  
^C  
--- 192.168.60.5 ping statistics ---  
13 packets transmitted, 0 received, 100% packet loss, time 12265ms  
  
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$
```

Wireshark:



2. Outside hosts can ping the router.
 - a. Command: ping seed-router (On Host A - 10.9.0.5)

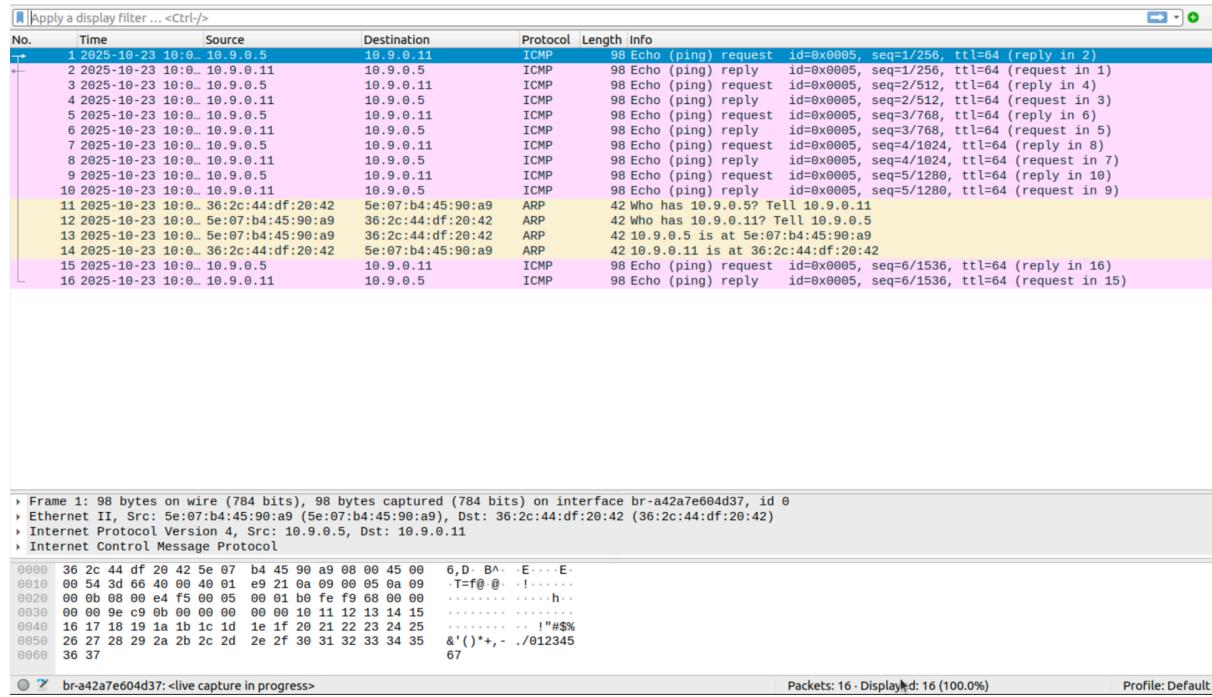
Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```

hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/
$ping seed-router
PING seed-router (10.9.0.11) 56(84) bytes of data.
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=1 ttl=64 time=0.43
ms
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=2 ttl=64 time=0.15
ms
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=3 ttl=64 time=0.14
ms
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=4 ttl=64 time=0.12
ms
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=5 ttl=64 time=0.14
ms
64 bytes from seed-router.net-10.9.0.0 (10.9.0.11): icmp_seq=6 ttl=64 time=0.09
ms
^C
--- seed-router ping statistics ---
6 packets transmitted, 6 received, 0% packet loss, time 5090ms
rtt min/avg/max/mdev = 0.097/0.182/0.430/0.112 ms
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/
$ 

```

Wireshark:



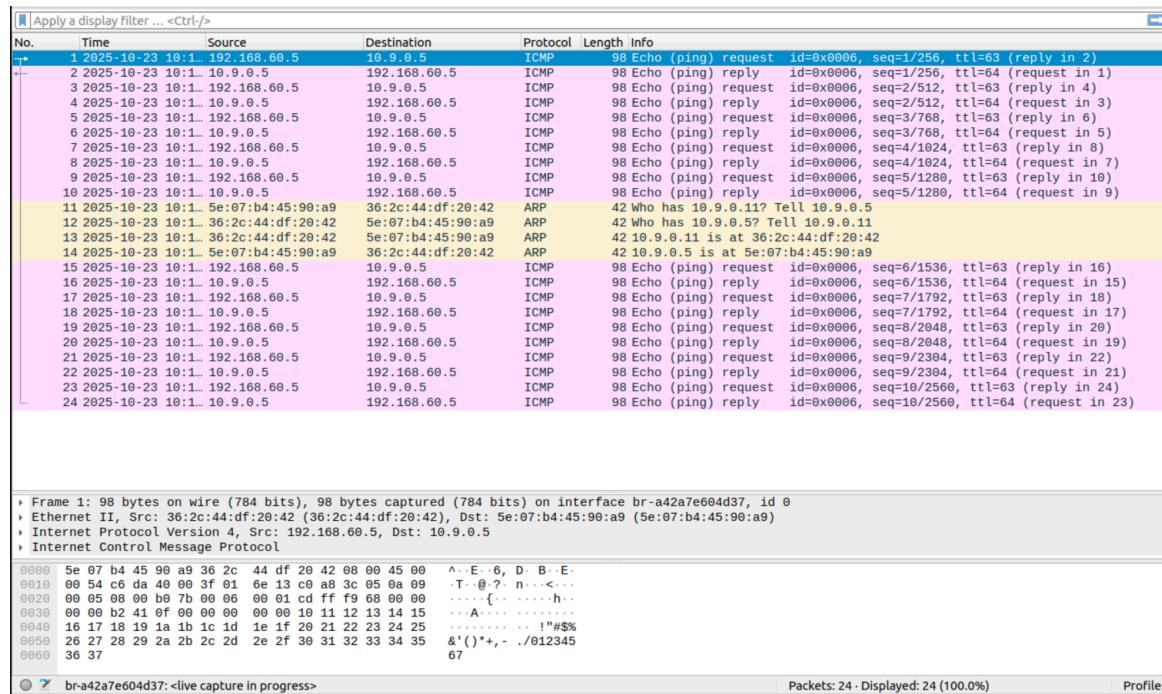
3. Internal hosts can ping Outside Hosts.

- a. Command: ping 10.9.0.5 (On host1-192.168.60.5)

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
host1-192.168.60.5:PES1UG23CS433:PranavHemanth/
$ping 10.9.0.5
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
64 bytes from 10.9.0.5: icmp_seq=1 ttl=63 time=1.27 ms
64 bytes from 10.9.0.5: icmp_seq=2 ttl=63 time=0.211 ms
64 bytes from 10.9.0.5: icmp_seq=3 ttl=63 time=0.163 ms
64 bytes from 10.9.0.5: icmp_seq=4 ttl=63 time=0.176 ms
64 bytes from 10.9.0.5: icmp_seq=5 ttl=63 time=0.229 ms
64 bytes from 10.9.0.5: icmp_seq=6 ttl=63 time=0.140 ms
64 bytes from 10.9.0.5: icmp_seq=7 ttl=63 time=0.116 ms
64 bytes from 10.9.0.5: icmp_seq=8 ttl=63 time=0.166 ms
64 bytes from 10.9.0.5: icmp_seq=9 ttl=63 time=0.178 ms
64 bytes from 10.9.0.5: icmp_seq=10 ttl=63 time=0.163 ms
^C
--- 10.9.0.5 ping statistics ---
10 packets transmitted, 10 received, 0% packet loss, time 9175ms
rtt min/avg/max/mdev = 0.116/0.281/1.269/0.330 ms
host1-192.168.60.5:PES1UG23CS433:PranavHemanth/
$
```

Wireshark:

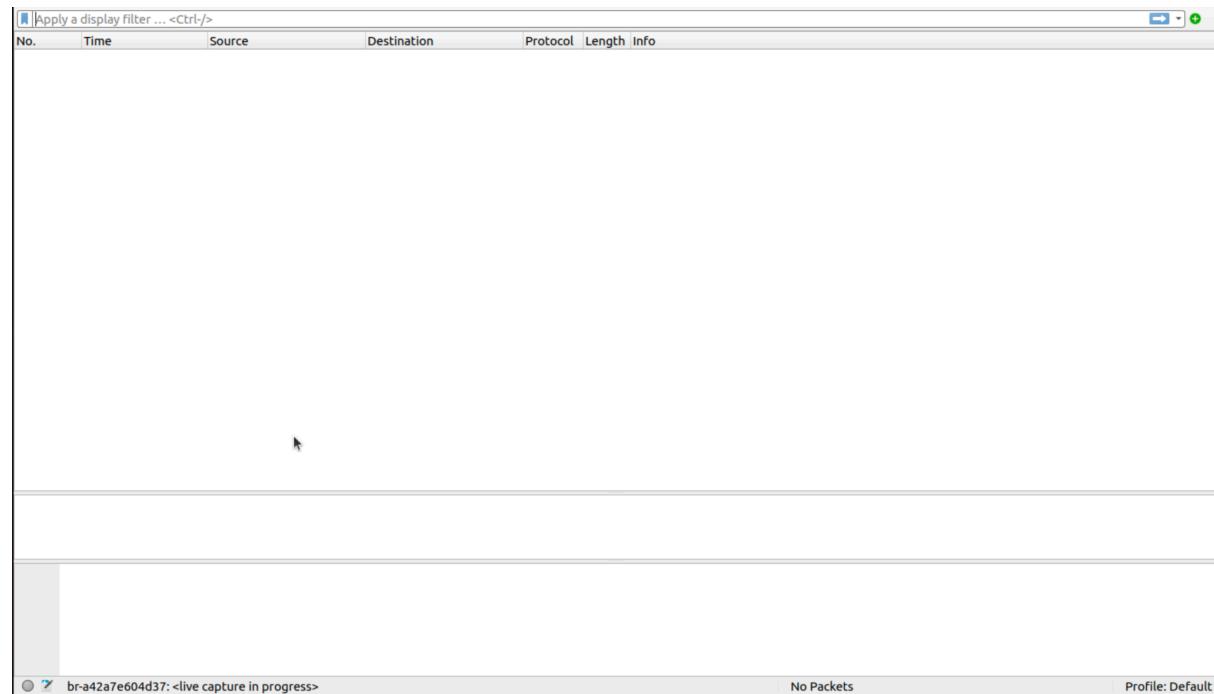


4. All other packets between the internal and external networks should be blocked.

- a. Command: telnet 10.9.0.5 (On host1-192.168.60.5)

```
host1-192.168.60.5:PES1UG23CS433:PranavHemanth/
$telnet 10.9.0.5
Trying 10.9.0.5...
^C
host1-192.168.60.5:PES1UG23CS433:PranavHemanth/
$
```

Wireshark:



(No packets show as they all are dropped)

Explanation:

In this task, we configured iptables on the seed-router to protect the internal network (192.168.60.0/24) while still allowing necessary connectivity through the FORWARD chain. We enforced a default-deny policy, only allowing specific ICMP traffic to be forwarded between the internal and external networks.

Three key rules were added:

- Block ICMP echo-requests coming from the external interface (eth0) to prevent outside hosts from pinging internal machines.
- Allow ICMP echo-requests originating from the internal network (eth1) so internal hosts can ping external systems.
- Allow ICMP echo-replies returning from external hosts so internal ping responses can be received.

With these rules active, the FORWARD chain policy was set to DROP, ensuring all other forwarded traffic is blocked.

Testing verified the expected behavior:

- Outside hosts could ping the router itself, since that uses the INPUT chain.
- Outside hosts could not ping internal hosts (blocked by our DROP rule).
- Internal hosts could ping external hosts because both request and reply ICMP packets were allowed.
- Other forwarded traffic, such as telnet, was blocked by the default DROP policy.

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

This experiment demonstrates how iptables can selectively restrict forwarding traffic to protect an internal network while still maintaining limited and controlled communication with external hosts.

Step3: Cleanup

25)Command:

```
iptables -F  
iptables -P OUTPUT ACCEPT  
iptables -P INPUT ACCEPT
```

```
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -F  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -P OUTPUT ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -P INPUT ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -t filter -L -n  
Chain INPUT (policy ACCEPT)  
target      prot opt source          destination  
  
Chain FORWARD (policy DROP)  
target      prot opt source          destination  
  
Chain OUTPUT (policy ACCEPT)  
target      prot opt source          destination  
seed-router:PES1UG23CS433:PranavHemanth:/  
$
```

Another way to restore the states of all the tables is to restart the container.

You can do it using the following command (you need to find the container's ID first):

```
$ docker restart <container ID> (Here - $ docker restart dccca3a58e38)
```

Task 2.C: Protecting Internal Servers

In this task, we want to protect the TCP servers inside the internal network (192.168.60.0/24).

More specifically, we would like to achieve the following objectives

1. All the internal hosts run a telnet server (listening to port 23). Outside hosts can only access the telnet server on 192.168.60.5, not the other internal hosts.
2. Outside hosts cannot access other internal servers.
3. Internal hosts can access all the internal servers.
4. Internal hosts cannot access external servers.

Step1: Execute the following iptables commands on the seed-router container-
26)Command:

```
iptables -A FORWARD -i eth0 -d 192.168.60.5 -p tcp --dport 23 -j ACCEPT  
iptables -A FORWARD -i eth1 -s 192.168.60.5 -p tcp --sport 23 -j ACCEPT  
iptables -P FORWARD DROP  
iptables -L -n -v
```

seed-router:

```
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -A FORWARD -i eth0 -d 192.168.60.5 -p tcp --dport 23 -j ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -A FORWARD -i eth1 -s 192.168.60.5 -p tcp --sport 23 -j ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -P FORWARD DROP  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -L -n -v  
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)  
 pkts bytes target     prot opt in     out      source          destination  
  
Chain FORWARD (policy DROP 0 packets, 0 bytes)  
 pkts bytes target     prot opt in     out      source          destination  
    0     0 ACCEPT      tcp   --  eth0    *       0.0.0.0/0        192.168.60.5      tcp dpt:23  
    0     0 ACCEPT      tcp   --  eth1    *       192.168.60.5     0.0.0.0/0      tcp spt:23  
  
Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)  
 pkts bytes target     prot opt in     out      source          destination  
seed-router:PES1UG23CS433:PranavHemanth:/  
$
```

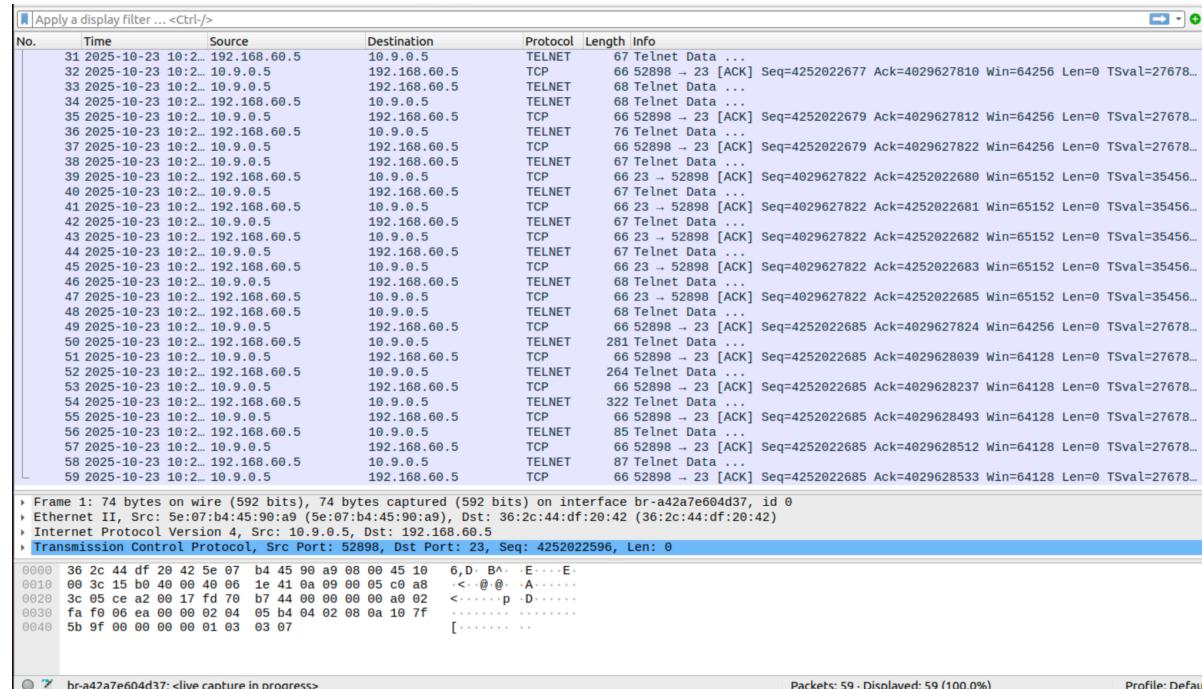
Step2: Now we shall see if these restrictions have been enforced in the network.

5. All the internal hosts run a telnet server (listening to port 23). Outside hosts can only access the telnet server on 192.168.60.5, not the other internal hosts.
 - a. Command: telnet 192.168.60.5 (On Host A - 10.9.0.5)

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$telnet 192.168.60.5  
Trying 192.168.60.5...  
Connected to 192.168.60.5.  
Escape character is '^]'.  
Ubuntu 20.04.6 LTS  
44e8f7e115a1 login: seed  
Password:  
Welcome to Ubuntu 20.04.6 LTS (GNU/Linux 5.15.0-157-generic aarch64)  
  
 * Documentation: https://help.ubuntu.com  
 * Management: https://landscape.canonical.com  
 * Support: https://ubuntu.com/advantage  
  
This system has been minimized by removing packages and content that are  
not required on a system that users do not log into.  
  
To restore this content, you can run the 'unminimize' command.  
  
The programs included with the Ubuntu system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/*copyright.  
  
Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by  
applicable law.  
seed@44e8f7e115a1:~$
```

Wireshark:



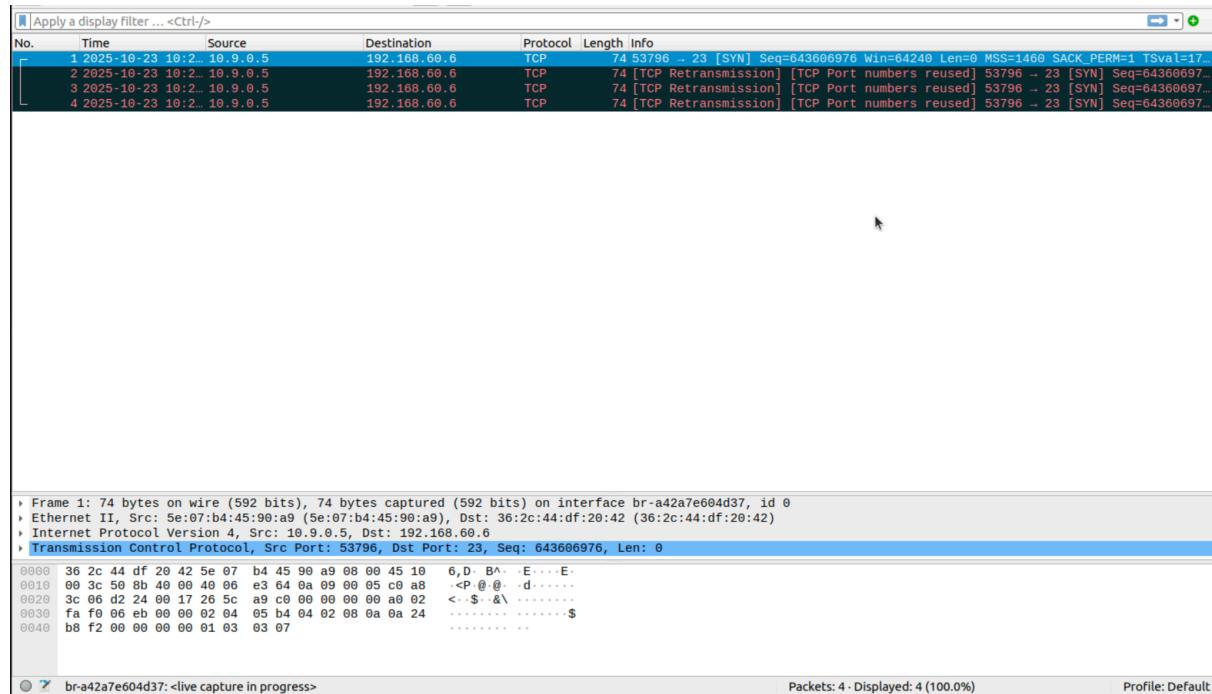
6. Outside hosts cannot access other internal servers.

- Command: telnet 192.168.60.6 (On Host A - 10.9.0.5)

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$telnet 192.168.60.6  
Trying 192.168.60.6...  
^C  
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$
```

Wireshark:

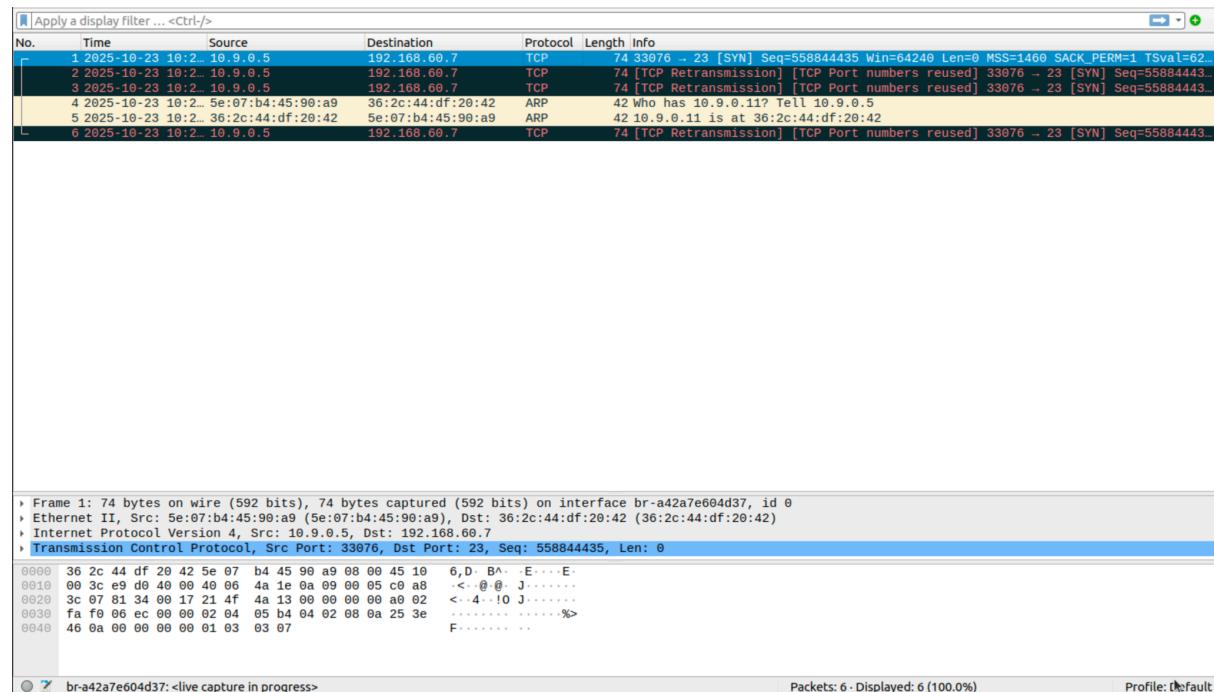


- b. Command: telnet 192.168.60.7 (On Host A - 10.9.0.5)

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$telnet 192.168.60.7  
Trying 192.168.60.7...  
^C  
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$
```

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

Wireshark:



7. Internal hosts can access all the internal servers.

- Command: telnet 192.168.60.5 (On host2 - 192.168.60.6)

```
host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/
$telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.6 LTS
44e8f7e115a1 login: seed
Password:
Welcome to Ubuntu 20.04.6 LTS (GNU/Linux 5.15.0-157-generic aarch64)

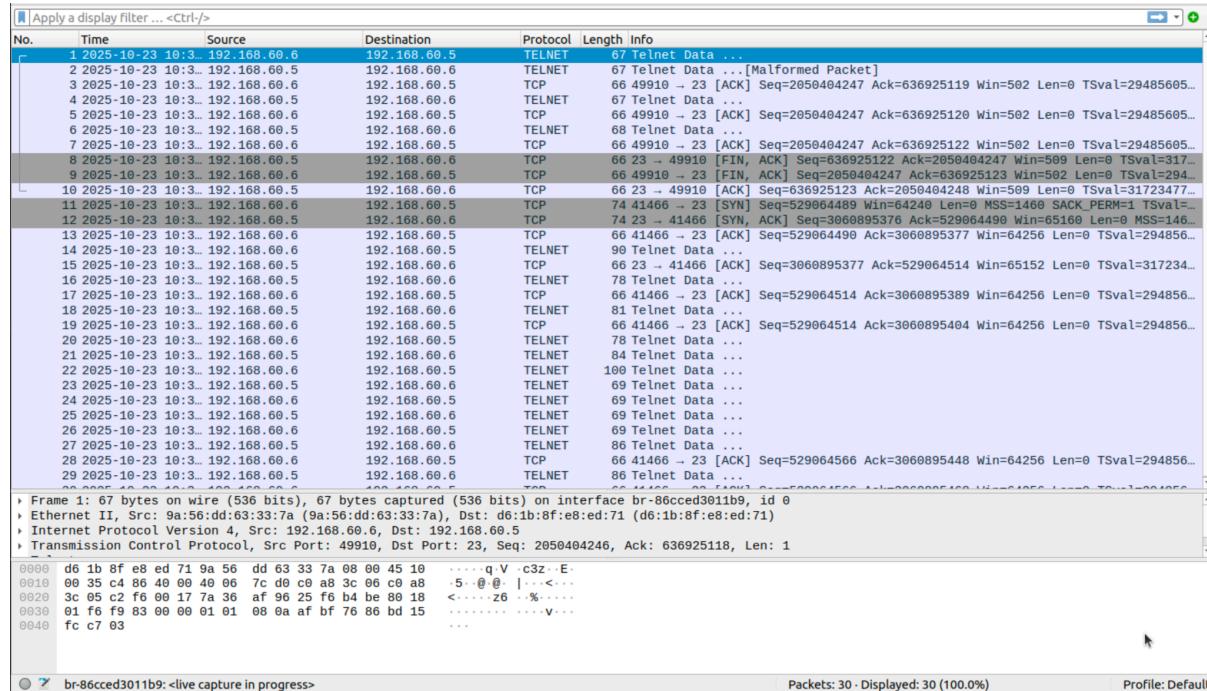
 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

This system has been minimized by removing packages and content that are
not required on a system that users do not log into.

To restore this content, you can run the 'unminimize' command.
Last login: Thu Oct 23 10:25:21 UTC 2025 from 10.9.0.5 on pts/2
seed@44e8f7e115a1:~$ █
```

Wireshark:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6



- b. Command: telnet 192.168.60.7 (On host2 - 192.168.60.6)

```

host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/
$telnet 192.168.60.7
Trying 192.168.60.7...
Connected to 192.168.60.7.
Escape character is '^]'.
Ubuntu 20.04.6 LTS
344d644c5360 login: seed
Password:
Welcome to Ubuntu 20.04.6 LTS (GNU/Linux 5.15.0-157-generic aarch64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

```

This system has been minimized by removing packages and content that are not required on a system that users do not log into.

To restore this content, you can run the 'unminimize' command.

The programs included with the Ubuntu system are free software; the exact distribution terms for each program are described in the individual files in /usr/share/doc/*/copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by applicable law.

```
seed@344d644c5360:~$
```

Wireshark:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

No.	Time	Source	Destination	Protocol	Length	Info
45	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TCP	66	23 -- 59942 [ACK] Seq=3487835586 Ack=1977414264 Win=65152 Len=0 TSval=42209...
46	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TELNET	67	Telnet Data ...
47	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TCP	66	23 -- 59942 [ACK] Seq=3487835586 Ack=1977414265 Win=65152 Len=0 TSval=42209...
48	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TELNET	67	Telnet Data ...
49	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TCP	66	23 -- 59942 [ACK] Seq=3487835586 Ack=1977414266 Win=65152 Len=0 TSval=42209...
50	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TELNET	68	Telnet Data ...
51	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TCP	66	23 -- 59942 [ACK] Seq=3487835586 Ack=1977414268 Win=65152 Len=0 TSval=42209...
52	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TELNET	68	Telnet Data ...
53	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TCP	66	59942 -- 23 [ACK] Seq=1977414268 Ack=3487835588 Win=64256 Len=0 TSval=49104...
54	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TELNET	357	Telnet Data ...
55	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TCP	66	59942 -- 23 [ACK] Seq=1977414268 Ack=3487835879 Win=64128 Len=0 TSval=49104...
56	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TELNET	188	Telnet Data ...
57	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TCP	66	59942 -- 23 [ACK] Seq=1977414268 Ack=3487836001 Win=64128 Len=0 TSval=49104...
58	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TELNET	341	Telnet Data ...
59	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TCP	66	59942 -- 23 [ACK] Seq=1977414268 Ack=3487836276 Win=64128 Len=0 TSval=49104...
60	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TELNET	87	Telnet Data ...
61	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TCP	66	59942 -- 23 [ACK] Seq=1977414268 Ack=3487836297 Win=64128 Len=0 TSval=49104...
62	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TELNET	75	Telnet Data ...
63	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TELNET	92	Telnet Data ...
64	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TCP	66	59942 -- 23 [ACK] Seq=1977414277 Ack=3487836323 Win=64128 Len=0 TSval=49104...
65	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TELNET	75	Telnet Data ...
66	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TELNET	92	Telnet Data ...
67	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TCP	66	59942 -- 23 [ACK] Seq=1977414286 Ack=3487836349 Win=64128 Len=0 TSval=49104...
68	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TELNET	75	Telnet Data ...
69	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TELNET	92	Telnet Data ...
70	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TCP	66	59942 -- 23 [ACK] Seq=1977414295 Ack=3487836375 Win=64128 Len=0 TSval=49105...
71	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TELNET	75	Telnet Data ...
72	2025-10-23 10:3...	192.168.60.7	192.168.60.6	TELNET	92	Telnet Data ...
73	2025-10-23 10:3...	192.168.60.6	192.168.60.7	TCP	66	59942 -- 23 [ACK] Seq=1977414304 Ack=3487836401 Win=64128 Len=0 TSval=49105...

▶ Frame 1: 74 bytes on wire (592 bits), 74 bytes captured (592 bits) on interface br-86cced3011b9, id 0
 ▶ Ethernet II, Src: 9a:56:dd:63:33:7a (9a:56:dd:63:33:7a), Dst: 00:0c:00:00:00:00 (00:0c:00:00:00:00)
 ▶ Internet Protocol Version 4, Src: 192.168.60.6, Dst: 192.168.60.7
 ▶ Transmission Control Protocol, Src Port: 59942, Dst Port: 23, Seq: 1977414177, Len: 0

```

0000 aa 2c 6d b0 dc 9a 56 dd 63 33 7a 08 00 45 10 >--.M-V .c3z-E.  

0010 00 3c f6 d7 40 00 46 06 a4 76 c0 a8 3c 00 c0 a8 <- @ @. Jv-<...  

0020 3c 07 ea 26 00 17 75 dc f2 21 00 00 00 a0 02 <- & u. !.....  

0030 fa f0 f9 8c 00 00 02 04 05 b4 04 02 08 0a 1d 44 .....@ .....D  

0040 af ea 00 00 00 00 01 03 03 07 .....  

  
```

br-86cced3011b9:<live capture in progress> Packets: 73 · Displayed: 73 (100.0%) Profile: Default

8. Internal hosts cannot access external servers.

- a. Command: telnet 10.9.0.5 (On host2 - 192.168.60.6)

```

host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/  

$telnet 10.9.0.5  

Trying 10.9.0.5...  

^C  

host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/  

$
```

Wireshark:

Apply a display filter ... <Ctrl-/>						
No.	Time	Source	Destination	Protocol	Length	Info
1	2025-10-23 10:3...	192.168.60.6	10.9.0.5	TCP	74	55180 -- 23 [SYN] Seq=69735237 Win=64240 Len=0 MSS=1460 SACK_PERM=1 TSval=230...
2	2025-10-23 10:3...	192.168.60.6	10.9.0.5	TCP	74	[TCP Retransmission] [TCP Port numbers reused] 55180 -- 23 [SYN] Seq=69735237...
3	2025-10-23 10:3...	192.168.60.6	10.9.0.5	TCP	74	[TCP Retransmission] [TCP Port numbers reused] 55180 -- 23 [SYN] Seq=69735237...
4	2025-10-23 10:3...	192.168.60.6	10.9.0.5	TCP	74	[TCP Retransmission] [TCP Port numbers reused] 55180 -- 23 [SYN] Seq=69735237...

▶ Frame 1: 74 bytes on wire (592 bits), 74 bytes captured (592 bits) on interface br-86cced3011b9, id 0
 ▶ Ethernet II, Src: 9a:56:dd:63:33:7a (9a:56:dd:63:33:7a), Dst: 3e:7e:a4:18:ce:cf (3e:7e:a4:18:ce:cf)
 ▶ Internet Protocol Version 4, Src: 192.168.60.6, Dst: 10.9.0.5
 ▶ Transmission Control Protocol, Src Port: 55180, Dst Port: 23, Seq: 69735237, Len: 0

```

0000 3e 7e a4 18 ce cf 9a 56 dd 63 33 7a 08 00 45 10 >--.M-V .c3z-E.  

0010 00 3c 91 74 40 00 46 06 a2 7b c0 a8 3c 00 c0 a8 <- @ @. {<...  

0020 00 05 d7 8c 00 17 04 28 13 45 00 00 00 00 a0 02 .....( E.....  

0030 fa f0 66 eb 00 00 02 04 05 b4 04 02 08 0a 89 5e .....A.....  

0040 20 df 00 00 00 00 01 03 03 07 .....  

  
```

br-86cced3011b9:<live capture in progress> Packets: 4 · Displayed: 4 (100.0%) Profile: Default

Explanation:

In this task, we configured the router's firewall using iptables to control access to internal TCP servers. The internal network (192.168.60.0/24) has multiple hosts running telnet servers, but only the server at 192.168.60.5 should be accessible from the external network. We added two specific ACCEPT rules to the FORWARD chain that allow only telnet traffic to and from 192.168.60.5, then set the default FORWARD policy to DROP so that all other forwarded traffic is blocked.

Testing confirmed the correct behavior:

- External hosts could successfully telnet into 192.168.60.5.
- External hosts could not access telnet servers on 192.168.60.6 or 192.168.60.7.
- Internal hosts could telnet to any internal server. These packets do not traverse the router and therefore bypass the FORWARD chain entirely.
- Internal hosts were blocked from telnetting to external systems, since those packets are forwarded through the router and hit the default DROP policy.

This demonstrates precise segmentation and protection of internal services. Access is granted only where explicitly permitted, while all other internal-external communications are restricted to prevent unauthorized entry or data exposure.

Step3: Cleanup

27)Command:

```
iptables -F  
iptables -P OUTPUT ACCEPT  
iptables -P INPUT ACCEPT
```

```
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -F  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -P OUTPUT ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -P INPUT ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -t filter -L -n  
chain INPUT (policy ACCEPT)  
target      prot opt source          destination  
  
chain FORWARD (policy DROP)  
target      prot opt source          destination  
  
chain OUTPUT (policy ACCEPT)  
target      prot opt source          destination  
seed-router:PES1UG23CS433:PranavHemanth:/  
$
```

Another way to restore the states of all the tables is to restart the container. You can do it using the following command (you need to find the container's ID first):
\$ docker restart <container ID> (Here - \$ docker restart dcca3a58e38)

Task 3: Connection Tracking and Stateful Firewall

In the previous task, we have only set up stateless firewalls, which inspect each packet independently. However, packets are usually not independent; they may be part of a TCP connection, or they may be ICMP packets triggered by other packets. Treating them independently does not take into consideration the context of the packets, and can thus lead to inaccurate, unsafe, or complicated firewall rules.

Task 3.A: Experiment with the Connection Tracking

To support stateful firewalls, we need to be able to track connections. This is achieved by the conntrack mechanism inside the kernel. In this task, we will conduct experiments related to this module, and get familiar with the connection tracking mechanism. In our experiment, we will check the connection tracking information on the router container.

This can be done using the following command:

```
# conntrack -L
```

The goal of the task is to use a series of experiments to help students understand the connection concept in this tracking mechanism, especially for the ICMP and UDP protocols, because unlike TCP , they do not have connections.

Step1: ICMP Experiment

28)Command: ping 192.168.60.5 (On Host A - 10.9.0.5)

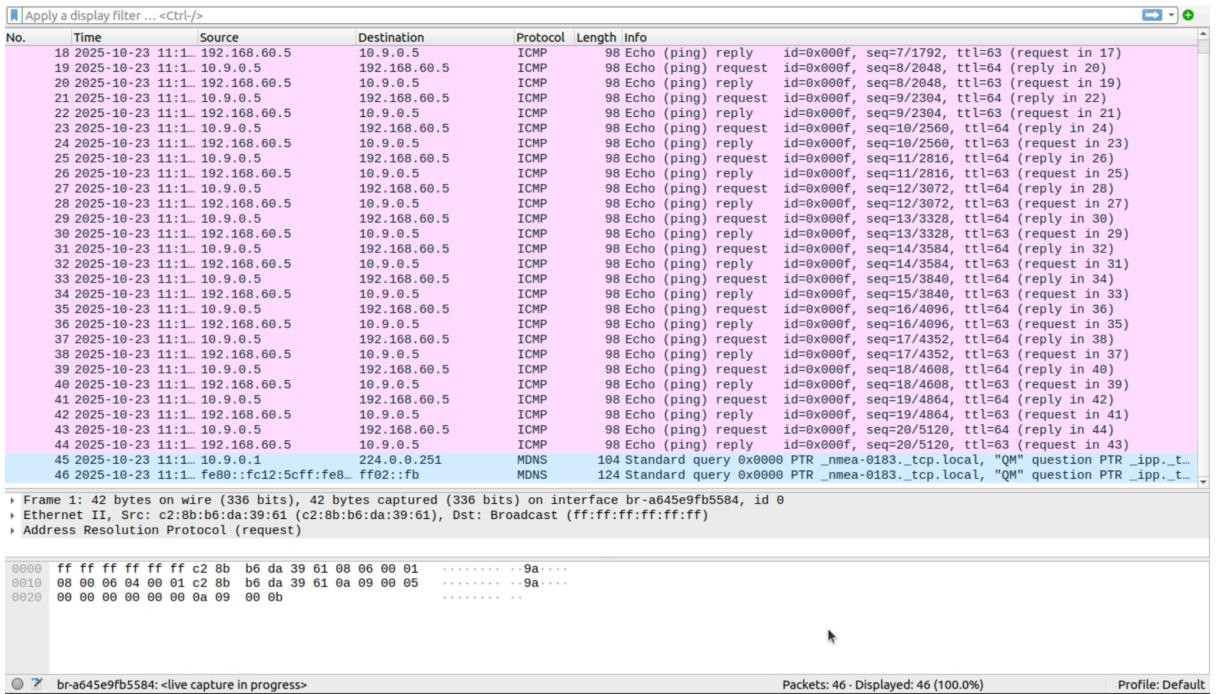
Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```

hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/ 
$ ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.831 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.160 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.153 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.143 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.233 ms
64 bytes from 192.168.60.5: icmp_seq=6 ttl=63 time=0.138 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.139 ms
64 bytes from 192.168.60.5: icmp_seq=8 ttl=63 time=0.150 ms
64 bytes from 192.168.60.5: icmp_seq=9 ttl=63 time=0.154 ms
64 bytes from 192.168.60.5: icmp_seq=10 ttl=63 time=0.154 ms
64 bytes from 192.168.60.5: icmp_seq=11 ttl=63 time=0.151 ms
64 bytes from 192.168.60.5: icmp_seq=12 ttl=63 time=0.144 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.155 ms
64 bytes from 192.168.60.5: icmp_seq=14 ttl=63 time=0.166 ms
64 bytes from 192.168.60.5: icmp_seq=15 ttl=63 time=0.081 ms
^C
--- 192.168.60.5 ping statistics ---
15 packets transmitted, 15 received, 0% packet loss, time 14325ms
rtt min/avg/max/mdev = 0.081/0.196/0.831/0.171 ms
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/ 
$ 

```

Wireshark:



29)Command: conntrack -L(On seed-router)

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
icmp    1 29 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=16 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=16 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
icmp    1 29 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=16 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=16 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
icmp    1 29 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=16 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=16 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
icmp    1 27 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=16 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=16 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
icmp    1 24 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=16 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=16 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
icmp    1 22 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=16 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=16 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
icmp    1 21 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=16 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=16 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
icmp    1 20 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=16 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=16 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
icmp    1 17 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=16 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=16 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

Questions:

How long can the ICMP connection state be kept?

Answer: ICMP entries remain in the conntrack table for **around 30 seconds** after the ping stops. The timeout counts down until it reaches zero, then the flow is removed.

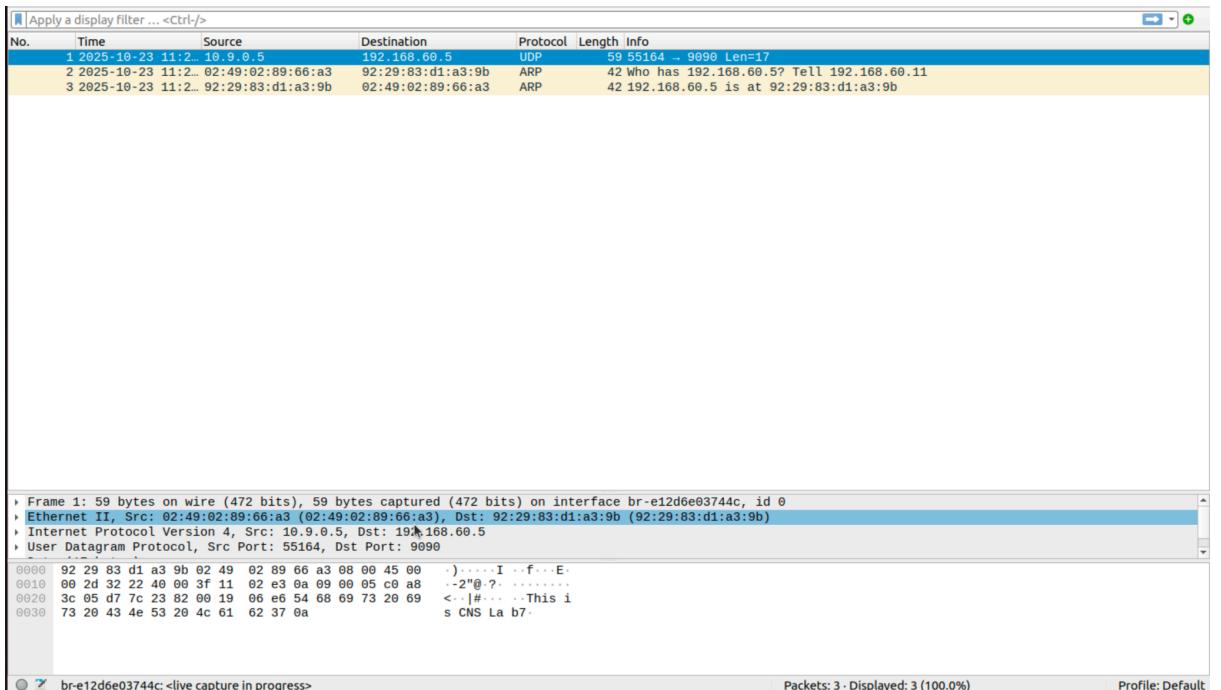
Step2: UDP Experiment

30)Command: nc -lu 9090 (On host1 - 192.168.60.5)

```
host1-192.168.60.5:PES1UG23CS433:PranavHemanth:/
$nc -lu 9090
This is CNS Lab7
[ ]
```

Wireshark:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6



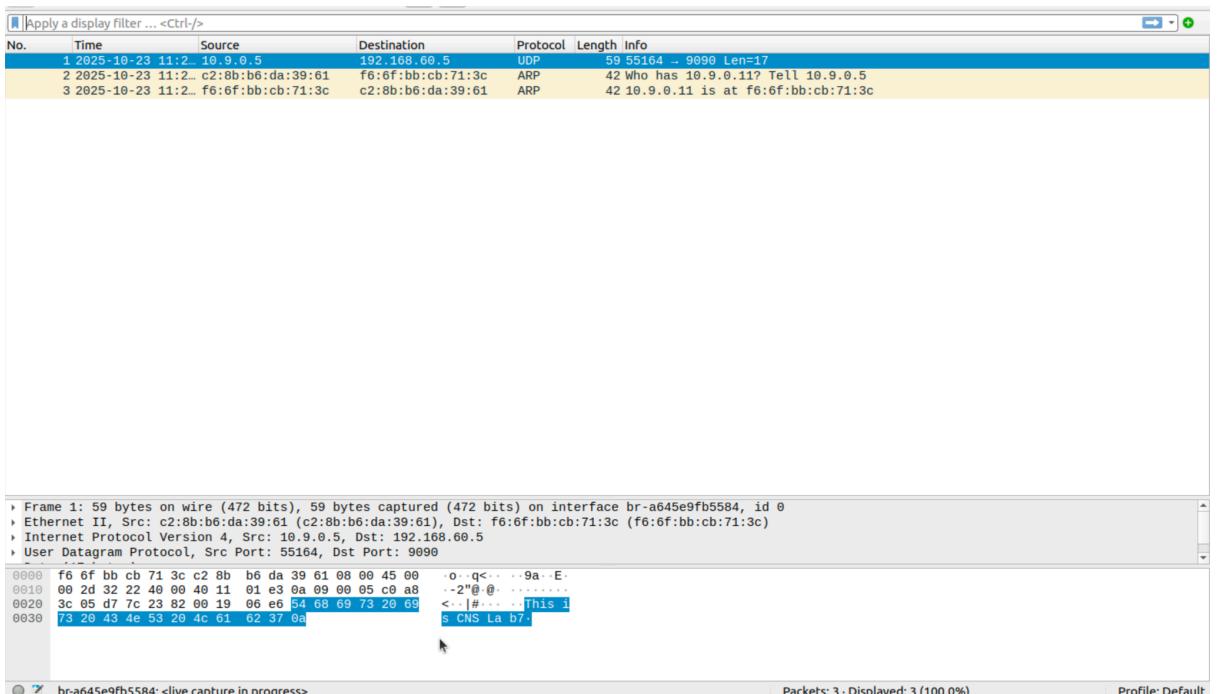
31)Command: nc -u 192.168.60.5 9090 (On Host A - 10.9.0.5)

```

hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/ 
$nc -u 192.168.60.5 9090
This is CNS Lab7

```

Wireshark:



32)Command: conntrack -L(On seed-router)

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
$conntrack -L
udp    17 15 src=10.9.0.5 dst=192.168.60.5 sport=50139 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=50139 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
udp    17 15 src=10.9.0.5 dst=192.168.60.5 sport=50139 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=50139 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
udp    17 14 src=10.9.0.5 dst=192.168.60.5 sport=50139 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=50139 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
udp    17 13 src=10.9.0.5 dst=192.168.60.5 sport=50139 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=50139 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
udp    17 7 src=10.9.0.5 dst=192.168.60.5 sport=50139 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=50139 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
udp    17 6 src=10.9.0.5 dst=192.168.60.5 sport=50139 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=50139 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
udp    17 3 src=10.9.0.5 dst=192.168.60.5 sport=50139 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=50139 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
udp    17 2 src=10.9.0.5 dst=192.168.60.5 sport=50139 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=50139 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$conntrack -L
conntrack v1.4.5 (conntrack-tools): 0 flow entries have been shown.
seed-router:PES1UG23CS433:PranavHemanth:/
$
```

Questions:

How long can the UDP connection state be kept?

Answer: UDP states are kept for a shorter period, roughly 10–30 seconds, and entries are marked UNREPLIED when there is no return packet. They disappear quickly due to the lack of session control or teardown.

Step3: TCP Experiment

33)Command: nc -l 9090 (On host1 - 192.168.60.5)

```
host1-192.168.60.5:PES1UG23CS433:PranavHemanth:/
$nc -l 9090
This is CNS Lab7
□
```

Wireshark:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

Apply a display filter ... <Ctrl-/>						
No.	Time	Source	Destination	Protocol	Length	Info
1	2025-10-23 11:3...	10.9.0.5	192.168.60.5	TCP	74	45754 → 9090 [SYN] Seq=1838322070 Win=64240 Len=0 MSS=1460 SACK_PERM=1 TStamp=18941562
2	2025-10-23 11:3...	192.168.60.5	10.9.0.5	TCP	74	9090 → 45754 [SYN, ACK] Seq=1838322671 Win=65160 Len=0 MSS=1460
3	2025-10-23 11:3...	10.9.0.5	192.168.60.5	TCP	66	45754 → 9090 [ACK] Seq=1838322671 Ack=11941563 Win=64256 Len=0 TStamp=2809110...
4	2025-10-23 11:3...	f6:6f:bb:cb:71:3c	c2:8b:b6:da:39:61	ARP	42	Who has 10.9.0.5? Tell 10.9.0.11
5	2025-10-23 11:3...	c2:8b:b6:da:39:61	f6:6f:bb:cb:71:3c	ARP	42	Who has 10.9.0.11? Tell 10.9.0.5
6	2025-10-23 11:3...	c2:8b:b6:da:39:61	f6:6f:bb:cb:71:3c	ARP	42	10.9.0.5 is at c2:8b:b6:da:39:61
7	2025-10-23 11:3...	f6:6f:bb:cb:71:3c	c2:8b:b6:da:39:61	ARP	42	10.9.0.11 is at f6:6f:bb:cb:71:3c
8	2025-10-23 11:3...	10.9.0.5	192.168.60.5	TCP	83	45754 → 9090 [PSH, ACK] Seq=1838322671 Ack=11941563 Win=64256 Len=17 TStamp=2...
9	2025-10-23 11:3...	192.168.60.5	10.9.0.5	TCP	66	9090 → 45754 [ACK] Seq=11941563 Ack=1838322688 Win=65152 Len=0 TStamp=3549818...

> Frame 1: 74 bytes on wire (592 bits), 74 bytes captured (592 bits) on interface br-a645e9fb5584, id 0
 > Ethernet II, Src: C: (c2:8b:b6:da:39:61) (f6:6f:bb:cb:71:3c) (f6:6f:bb:cb:71:3c)
 > Internet Protocol Version 4, Src: 10.9.0.5, Dst: 192.168.60.5
 > Transmission Control Protocol, Src Port: 45754, Dst Port: 9090, Seq: 1838322670, Len: 0

```
0000  f6 6f bb cb 71 3c c2 8b b6 da 39 61 08 00 45 00  ·o .q<· ·9a· E·
0010  00 3c e7 28 40 00 40 04 4c d8 0a 09 00 05 c0 a8  <- (@ ?) L·····
0020  3c 05 b2 ba 23 82 6d 92 93 ee 00 00 00 00 a0 02  <··#m······
0030  fa f0 06 ea 00 00 02 04 05 b4 04 02 08 0a 10 be  \····· ..
0040  5c f1 00 00 00 00 01 03 03 07  \····· ..
```

br-a645e9fb5584:<live capture in progress> Packets: 9 - Displayed: 9 (100.0%) Profile: Default

34)Command: nc 192.168.60.5 9090 (On Host A - 10.9.0.5)

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/ 
$nc 192.168.60.5 9090
This is CNS Lab7
```

Wireshark:

Apply a display filter ... <Ctrl-/>						
No.	Time	Source	Destination	Protocol	Length	Info
1	2025-10-23 11:3...	10.9.0.5	192.168.60.5	TCP	74	45754 → 9090 [SYN] Seq=1838322670 Win=64240 Len=0 MSS=1460 SACK_PERM=1 TStamp=18941562
2	2025-10-23 11:3...	192.168.60.5	10.9.0.5	TCP	74	9090 → 45754 [SYN, ACK] Seq=1838322671 Win=65160 Len=0 MSS=1460
3	2025-10-23 11:3...	10.9.0.5	192.168.60.5	TCP	66	45754 → 9090 [ACK] Seq=1838322671 Ack=11941563 Win=64256 Len=0 TStamp=2809110...
4	2025-10-23 11:3...	92:29:83:d1:a3:9b	02:49:02:89:66:a3	ARP	42	Who has 192.168.60.11? Tell 192.168.60.5
5	2025-10-23 11:3...	02:49:02:89:66:a3	92:29:83:d1:a3:9b	ARP	42	Who has 192.168.60.5? Tell 192.168.60.11
6	2025-10-23 11:3...	02:49:02:89:66:a3	92:29:83:d1:a3:9b	ARP	42	192.168.60.11 is at 02:49:02:89:66:a3
7	2025-10-23 11:3...	92:29:83:d1:a3:9b	02:49:02:89:66:a3	ARP	42	192.168.60.5 is at 92:29:83:d1:a3:9b
8	2025-10-23 11:3...	10.9.0.5	192.168.60.5	TCP	83	45754 → 9090 [PSH, ACK] Seq=1838322671 Ack=11941563 Win=64256 Len=17 TStamp=2...
9	2025-10-23 11:3...	192.168.60.5	10.9.0.5	TCP	66	9090 → 45754 [ACK] Seq=11941563 Ack=1838322688 Win=65152 Len=0 TStamp=3549818...

> Frame 1: 74 bytes on wire (592 bits), 74 bytes captured (592 bits) on interface br-e12d6e03744c, id 0
 > Ethernet II, Src: 02:49:02:89:66:a3 (02:49:02:89:66:a3), Dst: 92:29:83:d1:a3:9b (92:29:83:d1:a3:9b)
 > Internet Protocol Version 4, Src: 10.9.0.5, Dst: 192.168.60.5
 > Transmission Control Protocol, Src Port: 45754, Dst Port: 9090, Seq: 1838322670, Len: 0

```
0000  92 29 83 d1 a3 9b 02 49 02 89 66 a3 08 00 45 00  ·)···I ·f··E·
0010  00 3c e7 28 40 00 3f 06 4d d8 0a 09 00 05 c0 a8  <- (@ ?) M·····
0020  3c 05 b2 ba 23 82 6d 92 93 ee 00 00 00 00 a0 02  <··#m······
0030  fa f0 06 ea 00 00 02 04 05 b4 04 02 08 0a 10 be  \····· ..
0040  5c f1 00 00 00 00 01 03 03 07  \····· ..
```

br-e12d6e03744c:<live capture in progress> Packets: 9 - Displayed: 9 (100.0%) Profile: Default

35)Command: conntrack -L(On seed-router)

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
seed-router:PES1UG23CS433:PranavHemanth:/  
$conntrack -L  
tcp      6 431996 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=45754 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=45754 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.  
seed-router:PES1UG23CS433:PranavHemanth:/  
$conntrack -L  
tcp      6 431995 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=45754 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=45754 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.  
seed-router:PES1UG23CS433:PranavHemanth:/  
$conntrack -L  
tcp      6 431994 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=45754 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=45754 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.  
seed-router:PES1UG23CS433:PranavHemanth:/  
$conntrack -L  
tcp      6 431992 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=45754 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=45754 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.  
seed-router:PES1UG23CS433:PranavHemanth:/  
$conntrack -L  
tcp      6 431991 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=45754 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=45754 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.  
seed-router:PES1UG23CS433:PranavHemanth:/  
$conntrack -L  
tcp      6 431990 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=45754 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=45754 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.  
seed-router:PES1UG23CS433:PranavHemanth:/  
$conntrack -L  
tcp      6 431989 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=45754 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=45754 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.  
seed-router:PES1UG23CS433:PranavHemanth:/  
$conntrack -L  
tcp      6 431988 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=45754 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=45754 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.  
seed-router:PES1UG23CS433:PranavHemanth:/  
$conntrack -L  
tcp      6 431917 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=45754 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=45754 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.  
seed-router:PES1UG23CS433:PranavHemanth:/  
$conntrack -L  
tcp      6 431916 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=45754 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=45754 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.  
seed-router:PES1UG23CS433:PranavHemanth:/  
$conntrack -L  
tcp      6 431915 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=45754 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=45754 [ASSURED] mark=0 use=1  
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

Questions:

1. How long can the TCP connection state be kept?

Answer: TCP ESTABLISHED flows show a large timeout value of hundreds of thousands of seconds, typically about 5 days of idle timeout, because TCP has reliable session state and teardown behavior.

2. Close the TCP connection and execute conntrack -L on the router container.

Do you spot any difference?

Answer: Yes. Once the connection is closed:

- The state changes from ESTABLISHED to TIME_WAIT or CLOSE,
- Then the entry remains only briefly before being removed from the table.

This demonstrates proper tracking of TCP session lifecycle.

Explanation:

This task introduced connection tracking, which is essential for building stateful firewalls. Unlike stateless filtering, which treats every packet independently, conntrack keeps track of the state of communication between hosts. Even for protocols like ICMP and UDP that do not have a formal “connection,” conntrack still groups packets into flows and assigns temporary state.

Using the command conntrack -L on the seed-router, we inspected how the kernel tracks ICMP, UDP, and TCP traffic as we generated it from different hosts.

We observed:

- ICMP (ping) creates temporary tracked entries for echo-request and echo-reply packets, which expire shortly after traffic stops.
- UDP packets are also tracked briefly even if there is no reply. These entries disappear quickly since UDP lacks connection teardown.
- TCP has a well-defined state machine, so conntrack recognizes SYN, handshake completion, data transfer, and teardown. It keeps ESTABLISHED connections for a long timeout.

This helped show why stateful firewalls can allow return traffic without requiring explicit reverse-direction rules. The firewall can track and match packets as part of an already-approved flow, simplifying rule creation and improving security.

Task 3.B: Setting Up a Stateful Firewall

Now we are ready to set up firewall rules based on connections. In the following example, the "-m conntrack" option indicates that we are using the conntrack module, which is a very important module for iptables; it tracks connections, and iptables replies on the tracking information to build stateful firewalls. The --ctstate ESTABLISHED,RELATED indicates whether a packet belongs to an ESTABLISHED or RELATED connection. The rule allows TCP packets belonging to an existing connection to pass through.

For this task we have to rewrite the firewall rules in Task 2.C, but this time, we will add a rule allowing internal hosts to visit any external server (this was not allowed in Task 2.C).

Step1: Execute the following iptables commands on the seed-router container-
36)Command:

```
iptables -A FORWARD -p tcp -i eth0 -d 192.168.60.5 --dport 23 --syn -m conntrack --ctstate NEW -j ACCEPT
iptables -A FORWARD -i eth1 -p tcp --syn -m conntrack --ctstate NEW -j ACCEPT
iptables -A FORWARD -p tcp -m conntrack --ctstate RELATED,ESTABLISHED -j ACCEPT
iptables -A FORWARD -p tcp -j DROP
iptables -P FORWARD ACCEPT
iptables -L -n -v
```

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

seed-router:

```
seed-router:PES1UG23CS433:PranavHemanth:/
$>iptables -A FORWARD -p tcp -i eth0 -d 192.168.60.5 --dport 23 --syn -m conntrack --ctstate NEW -j ACCEPT
seed-router:PES1UG23CS433:PranavHemanth:/
$>iptables -A FORWARD -i eth1 -p tcp --syn -m conntrack --ctstate NEW -j ACCEPT
seed-router:PES1UG23CS433:PranavHemanth:/
$>iptables -A FORWARD -p tcp -m conntrack --ctstate RELATED,ESTABLISHED -j ACCEPT
seed-router:PES1UG23CS433:PranavHemanth:/
$>iptables -A FORWARD -p tcp -j DROP
seed-router:PES1UG23CS433:PranavHemanth:/
$>iptables -P FORWARD ACCEPT
seed-router:PES1UG23CS433:PranavHemanth:/
$>iptables -L -n -v
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source          destination
Chain FORWARD (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source          destination
  0    0 ACCEPT      tcp  --  eth0   *      0.0.0.0/0      192.168.60.5      tcp dpt:23 flags:0x17/0x02 ctstate NEW
  0    0 ACCEPT      tcp  --  eth1   *      0.0.0.0/0      0.0.0.0/0       tcp flags:0x17/0x02 ctstate NEW
  0    0 ACCEPT      tcp  --  *      *      0.0.0.0/0      0.0.0.0/0       ctstate RELATED,ESTABLISHED
  0    0 DROP        tcp  --  *      *      0.0.0.0/0      0.0.0.0/0
Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source          destination
seed-router:PES1UG23CS433:PranavHemanth:/
$>
```

Step2: Now we shall see if these restrictions have been enforced in the network.

9. All the internal hosts run a telnet server (listening to port 23). Outside hosts can only access the telnet server on 192.168.60.5, not the other internal hosts.

- a. Command: telnet 192.168.60.5 (On Host A - 10.9.0.5)

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/
$>telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
c04c5edada49 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.15.0-139-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:     https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

This system has been minimized by removing packages and content that are
not required on a system that users do not log into.

To restore this content, you can run the 'unminimize' command.

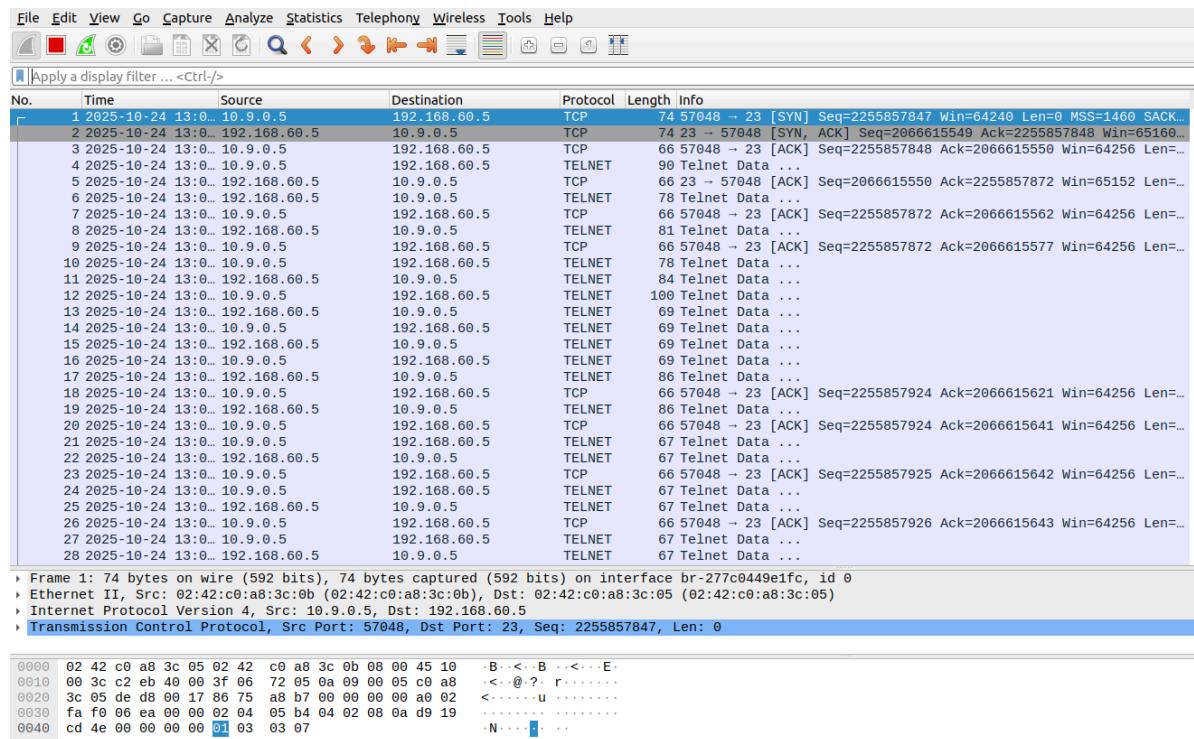
The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

seed@c04c5edada49:~$
```

Wireshark:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6



10. Outside hosts cannot access other internal servers.

- Command: telnet 192.168.60.6 (On Host A - 10.9.0.5)

```

hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/ 
$>telnet 192.168.60.6
Trying 192.168.60.6...
telnet: Unable to connect to remote host: Connection timed out
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/ 
$>

```

Wireshark:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

No.	Time	Source	Destination	Protocol	Length	Info
1	2025-10-24 13:1...	10.9.0.5	192.168.60.6	TCP	74	34892 -> 23 [SYN] Seq=3932607321 Win=64240 Len=0 MSS=1460 SACK...
2	2025-10-24 13:1...	10.9.0.5	192.168.60.6	TCP	74	40322 -> 23 [SYN] Seq=2417312943 Win=64240 Len=0 MSS=1460 SACK...
3	2025-10-24 13:1...	10.9.0.5	192.168.60.6	TCP	74	[TCP Retransmission] 40322 -> 23 [SYN] Seq=2417312943 Win=6424...
4	2025-10-24 13:1...	10.9.0.5	192.168.60.6	TCP	74	[TCP Retransmission] 40322 -> 23 [SYN] Seq=2417312943 Win=6424...
5	2025-10-24 13:1...	10.9.0.5	192.168.60.6	TCP	74	[TCP Retransmission] 40322 -> 23 [SYN] Seq=2417312943 Win=6424...
6	2025-10-24 13:1...	02:42:0a:09:00:05	02:42:0a:09:00:0b	ARP	42	Who has 10.9.0.11? Tell 10.9.0.5
7	2025-10-24 13:1...	02:42:0a:09:00:05	02:42:0a:09:00:05	ARP	42	10.9.0.11 is at 02:42:0a:09:00:0b
8	2025-10-24 13:1...	10.9.0.5	192.168.60.6	TCP	74	[TCP Retransmission] 40322 -> 23 [SYN] Seq=2417312943 Win=6424...
9	2025-10-24 13:1...	10.9.0.5	192.168.60.6	TCP	74	[TCP Retransmission] 40322 -> 23 [SYN] Seq=2417312943 Win=6424...
10	2025-10-24 13:1...	10.9.0.5	192.168.60.6	TCP	74	[TCP Retransmission] 40322 -> 23 [SYN] Seq=2417312943 Win=6424...
11	2025-10-24 13:1...	02:42:0a:09:00:05	02:42:0a:09:00:0b	ARP	42	Who has 10.9.0.11? Tell 10.9.0.5
12	2025-10-24 13:1...	02:42:0a:09:00:05	02:42:0a:09:00:05	ARP	42	10.9.0.11 is at 02:42:0a:09:00:0b

Frame 1: 74 bytes on wire (592 bits), 74 bytes captured (592 bits) on interface br-e660c7a8e614, id 0
 > Ethernet II, Src: 02:42:0a:09:00:05 (02:42:0a:09:00:05), Dst: 02:42:0a:09:00:0b (02:42:0a:09:00:0b)
 > Internet Protocol Version 4, Src: 10.9.0.5, Dst: 192.168.60.6
 > Transmission Control Protocol, Src Port: 34892, Dst Port: 23, Seq: 3932607321, Len: 0

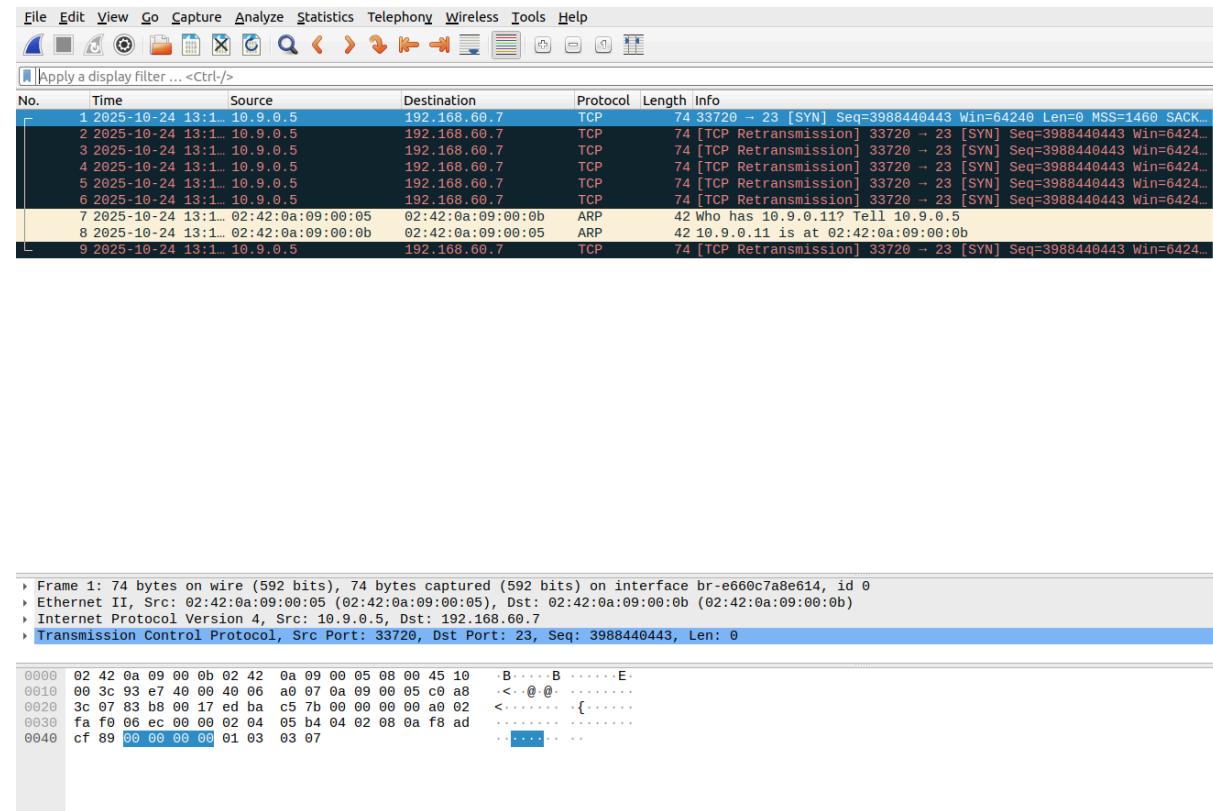
0000	02	42	0a	09	00	0b	02	42	0a	09	00	05	08	00	45	10	.B.	..BE.
0010	00	3c	e3	a6	49	00	40	06	50	49	00	09	00	05	c0	a8	<- .@.	PI
0020	3c	06	88	4c	00	17	ea	66	d3	59	00	00	00	00	ad	02	<- L.	f	Y.....
0030	fa	f0	6e	00	00	02	04	05	b4	04	02	08	00	a8	e5
0040	0f	38	00	00	00	00	01	03	03	07	8.	

b. Command: telnet 192.168.60.7 (On Host A - 10.9.0.5)

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/
$>telnet 192.168.60.7
Trying 192.168.60.7...
telnet: Unable to connect to remote host: Connection timed out
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/
$>
```

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

Wireshark:



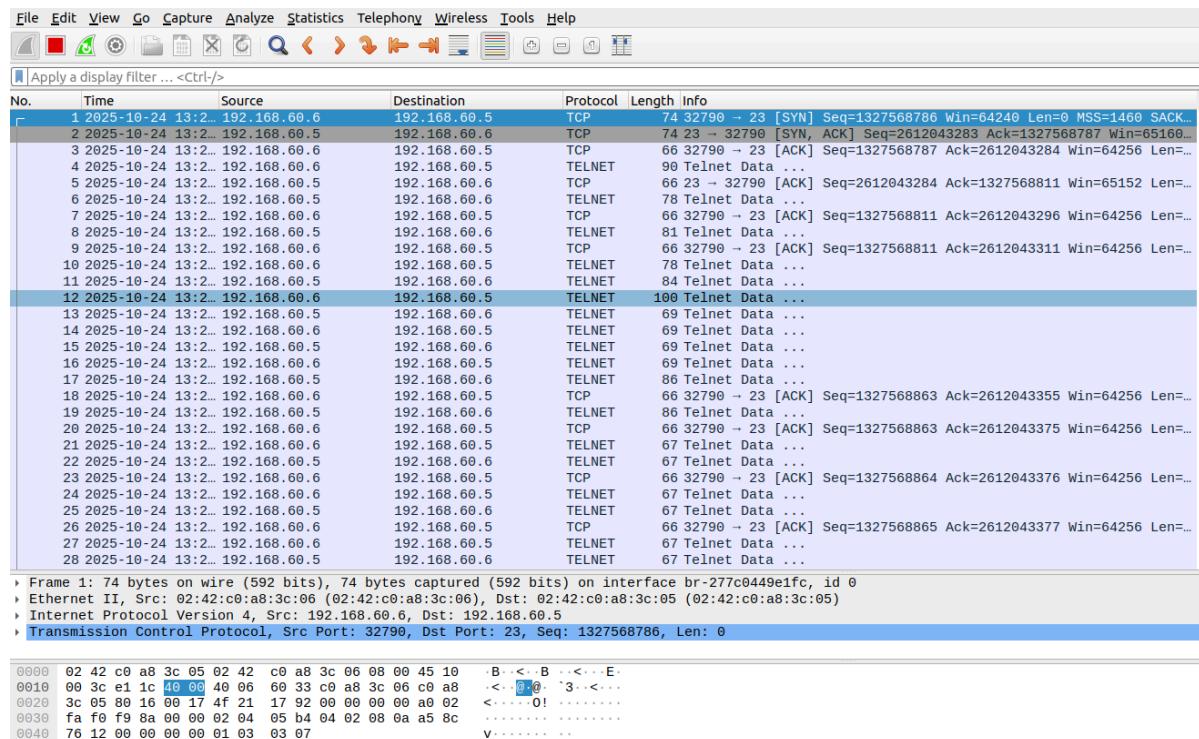
11. Internal hosts can access all the internal servers.

- Command: telnet 192.168.60.5 (On host2 - 192.168.60.6)

```
host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/  
$>telnet 192.168.60.5  
Trying 192.168.60.5...  
Connected to 192.168.60.5.  
Escape character is '^]'.  
Ubuntu 20.04.1 LTS  
c04c5edada49 login: seed  
Password:  
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.15.0-139-generic x86_64)  
  
* Documentation: https://help.ubuntu.com  
* Management: https://landscape.canonical.com  
* Support: https://ubuntu.com/advantage  
  
This system has been minimized by removing packages and content that are  
not required on a system that users do not log into.  
  
To restore this content, you can run the 'unminimize' command.  
Last login: Fri Oct 24 17:20:17 UTC 2025 from www.SeedLabSQLInjection.com on pts/2  
seed@c04c5edada49:~$
```

Wireshark:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6



- b. Command: telnet 192.168.60.7 (On host2 - 192.168.60.6)

```
host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/
$>telnet 192.168.60.7
Trying 192.168.60.7...
Connected to 192.168.60.7.
Escape character is '^].
Ubuntu 20.04.1 LTS
d6239420659d login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.15.0-139-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:      https://landscape.canonical.com
 * Support:         https://ubuntu.com/advantage
```

This system has been minimized by removing packages and content that are not required on a system that users do not log into.

To restore this content, you can run the 'unminimize' command.

The programs included with the Ubuntu system are free software; the exact distribution terms for each program are described in the individual files in /usr/share/doc/*/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by applicable law.

```
seed@d6239420659d:~$ █
```

Wireshark:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

No.	Time	Source	Destination	Protocol	Length	Info
1	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TCP	74	54852 → 23 [SYN] Seq=3128126267 Win=64240 Len=0 MSS=1460 SACK...
2	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TCP	74	23 → 54852 [SYN, ACK] Seq=3128126268 Ack=3128126268 Win=65160 Len=...
3	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TCP	66	54852 → 23 [ACK] Seq=3128126268 Ack=3128126268 Win=64256 Len=...
4	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TELNET	90	Telnet Data ...
5	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TCP	66	23 → 54852 [ACK] Seq=3128126292 Win=65152 Len=...
6	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TELNET	78	Telnet Data ...
7	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TCP	66	54852 → 23 [ACK] Seq=3128126292 Ack=3128126292 Win=64256 Len=...
8	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TELNET	81	Telnet Data ...
9	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TCP	66	54852 → 23 [ACK] Seq=3128126292 Ack=3128126292 Win=64256 Len=...
10	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TELNET	78	Telnet Data ...
11	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TELNET	84	Telnet Data ...
12	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TELNET	100	Telnet Data ...
13	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TELNET	69	Telnet Data ...
14	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TELNET	69	Telnet Data ...
15	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TELNET	69	Telnet Data ...
16	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TELNET	69	Telnet Data ...
17	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TELNET	86	Telnet Data ...
18	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TCP	66	54852 → 23 [ACK] Seq=3128126344 Ack=3128126344 Win=64256 Len=...
19	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TELNET	86	Telnet Data ...
20	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TCP	66	54852 → 23 [ACK] Seq=3128126344 Ack=3128126344 Win=64256 Len=...
21	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TELNET	67	Telnet Data ...
22	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TELNET	67	Telnet Data ...
23	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TCP	66	54852 → 23 [ACK] Seq=3128126345 Ack=3128126345 Win=64256 Len=...
24	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TELNET	67	Telnet Data ...
25	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TELNET	67	Telnet Data ...
26	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TCP	66	54852 → 23 [ACK] Seq=3128126346 Ack=3128126346 Win=64256 Len=...
27	2025-10-24 13:2...	192.168.60.6	192.168.60.7	TELNET	67	Telnet Data ...
28	2025-10-24 13:2...	192.168.60.7	192.168.60.6	TELNET	67	Telnet Data ...

Frame 1: 74 bytes on wire (592 bits), 74 bytes captured (592 bits) on interface br-277c0449e1fc, id 0
 Ethernet II, Src: 02:42:c0:a3:07 (02:42:c0:a3:07), Dst: 02:42:c0:a8:3c:07 (02:42:c0:a8:3c:07)
 Internet Protocol Version 4, Src: 192.168.60.6, Dst: 192.168.60.7
 Transmission Control Protocol, Src Port: 54852, Dst Port: 23, Seq: 3128126267, Len: 0

```

0000  02 42 c0 a3 07 02 42 c0 a8 3c 06 08 00 45 10  .B-<.B ..<..E
0010  00 3c 3e cd 09 00 40 06 02 81 c0 a8 3c 06 c0 a8  .<> @ @. ....<...
0020  3c 07 d6 44 00 17 ba 73 6b 3b 00 00 00 00 a0 02  .<-D--s k;.....
0030  fa f0 f9 8c 00 00 02 04 05 b4 04 02 08 0a 22 af  .....
0040  c1 04 00 00 00 00 01 03 03 07  ..... .

```

12. Internal hosts can access external servers.

- a. Command: telnet 10.9.0.5 (On host2 - 192.168.60.6)

```

host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/
$>telnet 10.9.0.5
Trying 10.9.0.5...
Connected to 10.9.0.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
972e43bfc613 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.15.0-139-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:      https://landscape.canonical.com
 * Support:         https://ubuntu.com/advantage

```

This system has been minimized by removing packages and content that are not required on a system that users do not log into.

To restore this content, you can run the 'unminimize' command.

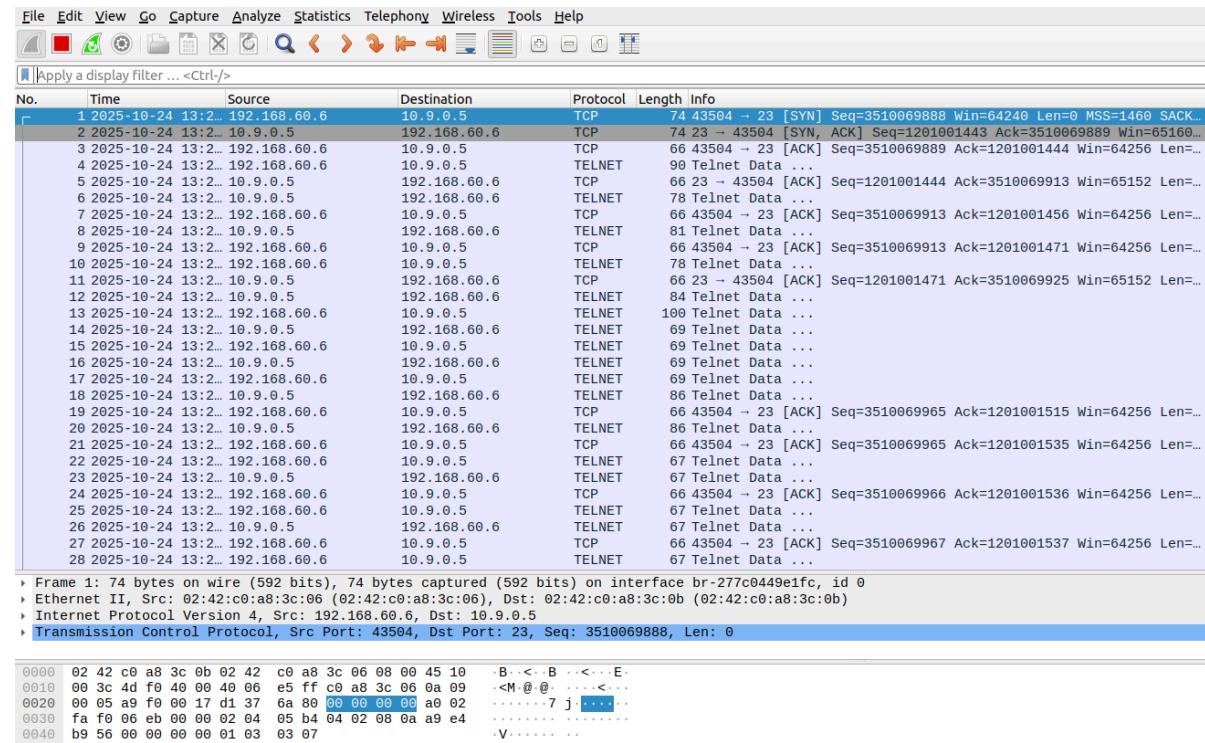
The programs included with the Ubuntu system are free software; the exact distribution terms for each program are described in the individual files in /usr/share/doc/*/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by applicable law.

seed@972e43bfc613:~\$

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

Wireshark:



Explanation:

In this task, we switched from a stateless firewall to a stateful one using the kernel's connection tracking module (conntrack). Stateful firewalls understand the context of network traffic. Instead of inspecting each packet independently, they track active connections and determine whether a packet belongs to a valid, ongoing session.

We rewrote the firewall rules so that:

- External hosts are only allowed to start Telnet connections to one internal server: 192.168.60.5:23
- Internal hosts can initiate any external TCP connection
- Return traffic for accepted connections is automatically allowed due to ESTABLISHED,RELATED handling
- All other TCP traffic is dropped

Once the rules were applied, we tested different connection scenarios:

- From Host A (external), Telnet to 192.168.60.5 succeeded because it matched the NEW allowed rule.
- External Telnet to 192.168.60.6 or .7 failed since no rule permits NEW connections to them.
- Communication between internal hosts always worked because it never touches the router's FORWARD chain.
- Internal hosts could now access external servers, unlike in Task 2.C, because NEW outbound connections are explicitly allowed.

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

Overall, connection tracking made our firewall smarter, simpler, and more secure by allowing replies for approved connections while blocking all unauthorized attempts.

The firewall rules for the stateful setup explained individually are as follows:

1. Allow NEW TCP SYN from the Internet to 192.168.60.5:23
 - This rule ensures that only the intended internal server is reachable from outside, restricting external access to other internal hosts.
2. Allow NEW TCP SYN from the internal LAN to any destination
 - This rule enables internal hosts to initiate connections to external servers, supporting normal outbound traffic.
3. Allow RELATED or ESTABLISHED connections
 - This rule permits reply traffic for approved connections. It forms the core of the stateful firewall, automatically allowing responses without additional rules, simplifying management and improving security.
4. Drop all remaining TCP traffic
 - This rule blocks unauthorized connection attempts and acts as a security cleanup measure, ensuring that any traffic not explicitly allowed is discarded.

Step3: Cleanup

37)Command:

```
iptables -F  
iptables -P OUTPUT ACCEPT  
iptables -P INPUT ACCEPT
```

```
seed-router:PES1UG23CS433:PranavHemanth:/  
$ iptables -F  
seed-router:PES1UG23CS433:PranavHemanth:/  
$ iptables -P OUTPUT ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$ iptables -P INPUT ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$ iptables -t filter -L -n  
Chain INPUT (policy ACCEPT)  
target     prot opt source          destination  
  
Chain FORWARD (policy DROP)  
target     prot opt source          destination  
  
Chain OUTPUT (policy ACCEPT)  
target     prot opt source          destination  
seed-router:PES1UG23CS433:PranavHemanth:/  
$
```

Another way to restore the states of all the tables is to restart the container.

You can do it using the following command (you need to find the container's ID first):

```
$ docker restart <container ID> (Here - $ docker restart dccca3a58e38)
```

Task 4: Limiting Network Traffic

In addition to blocking packets, we can also limit the number of packets that can pass through the firewall. This can be done using the limit module of iptables. In this task, we will use this module to limit how many packets from 10.9.0.5 are allowed to get into the internal network.

You can use "iptables -m limit -h" to see the manual Please run the following commands on the router, and then ping 192.168.60.5 from 10.9.0.5.

Conduct the experiment with and without the second rule, and then explain whether the second rule is needed or not, and why.

Step1: On seed router execute (With 2nd rule)

38)Command:

```
iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute --limit-burst 5 -j ACCEPT
iptables -A FORWARD -s 10.9.0.5 -j DROP
iptables -L -n -v
```

seed-router:

```
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute --limit-burst 5 -j ACCEPT
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -A FORWARD -s 10.9.0.5 -j DROP
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -L -n -v
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source               destination
Chain FORWARD (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source               destination
      0     0 ACCEPT     all  --  *      *      10.9.0.5          0.0.0.0/0           limit: avg 10/min burst 5
      0     0 DROP       all  --  *      *      10.9.0.5          0.0.0.0/0
Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source               destination
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -t filter -L -n
Chain INPUT (policy ACCEPT)
target     prot opt source
Chain FORWARD (policy ACCEPT)
target     prot opt source               destination
ACCEPT    all  --  10.9.0.5          0.0.0.0/0           limit: avg 10/min burst 5
DROP      all  --  10.9.0.5          0.0.0.0/0
Chain OUTPUT (policy ACCEPT)
target     prot opt source               destination
seed-router:PES1UG23CS433:PranavHemanth:/
$
```

Step2: Now we shall see if these restrictions have been enforced in the network.

39)Command: ping 192.168.60.5 (On Host A - 10.9.0.5)

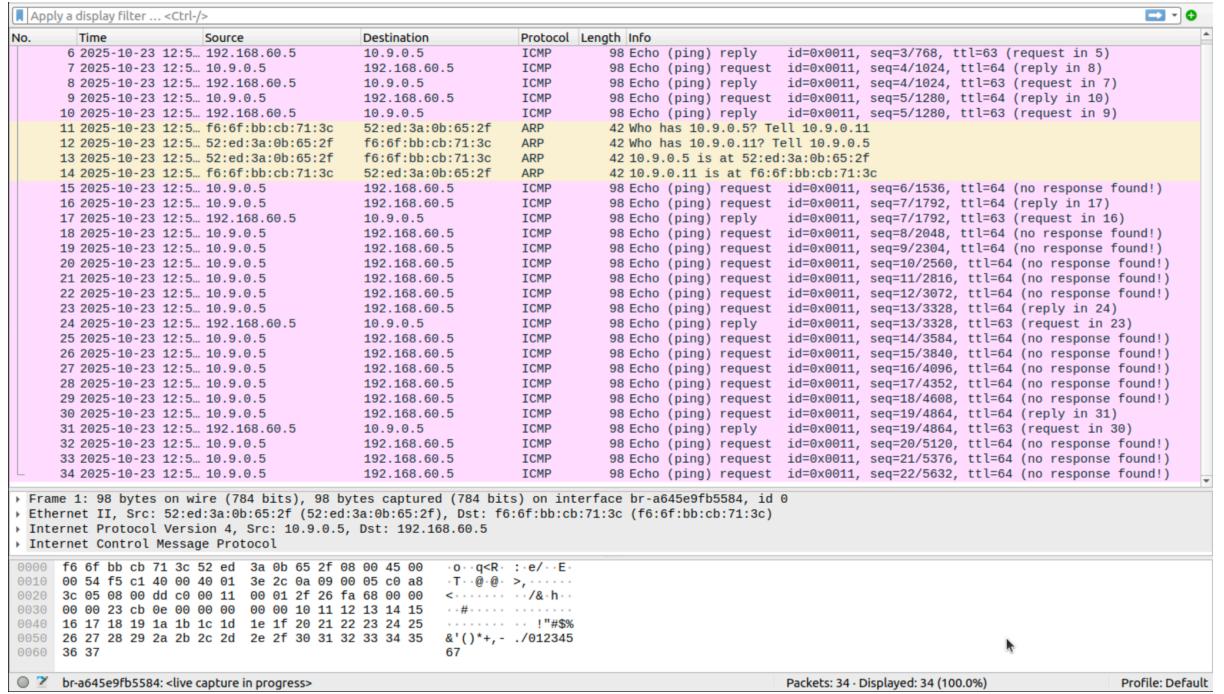
Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```

hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/ 
$ ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=2.25 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.263 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.224 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.192 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.228 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.141 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.231 ms
64 bytes from 192.168.60.5: icmp_seq=19 ttl=63 time=0.186 ms
^C
--- 192.168.60.5 ping statistics ---
22 packets transmitted, 8 received, 63.6364% packet loss, time 21489ms
rtt min/avg/max/mdev = 0.141/0.464/2.250/0.675 ms
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/ 
$ 

```

Wireshark:



Step3: Cleanup

40)Command:

iptables -F

iptables -P OUTPUT ACCEPT

iptables -P INPUT ACCEPT

```
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -F  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -P OUTPUT ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -P INPUT ACCEPT  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -t filter -L -n  
Chain INPUT (policy ACCEPT)  
target     prot opt source          destination  
  
Chain FORWARD (policy DROP)  
target     prot opt source          destination  
  
Chain OUTPUT (policy ACCEPT)  
target     prot opt source          destination  
seed-router:PES1UG23CS433:PranavHemanth:/  
$
```

Another way to restore the states of all the tables is to restart the container. You can do it using the following command (you need to find the container's ID first):
\$ docker restart <container ID> (Here - \$ docker restart dccca3a58e38)

Explanation: Experiment 1:- With the Second DROP Rule

In this experiment, the firewall on the router was configured with two FORWARD rules for traffic from 10.9.0.5. The first rule uses the limit module to accept up to 10 packets per minute with a burst of 5, and the second rule explicitly drops any remaining packets from the same source that exceed this limit.

Testing this setup by pinging 192.168.60.5 from Host A shows that the connection experiences significant packet loss (about 70%), which demonstrates that the firewall is correctly throttling traffic. The initial burst of packets is allowed through, followed by a slow trickle at the defined rate, while excess packets are discarded by the DROP rule. Without this second rule, the packets exceeding the limit would fall through to the chain's default policy and be accepted, rendering the rate-limiting ineffective.

Purpose of the rules:

1. iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute --limit-burst 5 -j ACCEPT — defines the allowance, accepting packets that are within the specified rate.
2. iptables -A FORWARD -s 10.9.0.5 -j DROP — enforces the limit by dropping packets that exceed the allowance.

Step4: On seed router execute (Without 2nd rule)

41)Command:

```
iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute --limit-burst 5 -j ACCEPT
iptables -L -n -v
```

seed-router:

```
seed-router:PES1UG23CS433:PranavHemanth:/
$ iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute --limit-burst 5 -j ACCEPT
$ iptables -L -n -v
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source          destination
Chain FORWARD (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source          destination
      0      0 ACCEPT     all  --  *       *      10.9.0.5          0.0.0.0/0
      limit: avg 10/min burst 5
Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target     prot opt in     out     source          destination
seed-router:PES1UG23CS433:PranavHemanth:/
$
```

Step5: Now we shall see if these restrictions have been enforced in the network.

42)Command: ping 192.168.60.5 (On Host A - 10.9.0.5)

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/
$ ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
 64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.557 ms
 64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.165 ms
 64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.226 ms
 64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.184 ms
 64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.209 ms
 64 bytes from 192.168.60.5: icmp_seq=6 ttl=63 time=0.196 ms
 64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.183 ms
 64 bytes from 192.168.60.5: icmp_seq=8 ttl=63 time=0.178 ms
 64 bytes from 192.168.60.5: icmp_seq=9 ttl=63 time=0.183 ms
 64 bytes from 192.168.60.5: icmp_seq=10 ttl=63 time=0.111 ms
 64 bytes from 192.168.60.5: icmp_seq=11 ttl=63 time=0.166 ms
 64 bytes from 192.168.60.5: icmp_seq=12 ttl=63 time=0.179 ms
^C
--- 192.168.60.5 ping statistics ---
12 packets transmitted, 12 received, 0% packet loss, time 11252ms
rtt min/avg/max/mdev = 0.111/0.211/0.557/0.107 ms
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/
$
```

43)

Wireshark:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

No.	Time	Source	Destination	Protocol	Length	Info
2	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) reply id=0x0012, seq=1/256, ttl=63 (request in 1)
3	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) request id=0x0012, seq=2/512, ttl=64 (reply in 4)
4	2025-10-23 15:2...	192.168.60.5	10.9.0.5	ICMP	98	Echo (ping) reply id=0x0012, seq=2/512, ttl=63 (request in 3)
5	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) request id=0x0012, seq=3/768, ttl=64 (reply in 6)
6	2025-10-23 15:2...	192.168.60.5	10.9.0.5	ICMP	98	Echo (ping) reply id=0x0012, seq=3/768, ttl=63 (request in 5)
7	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) request id=0x0012, seq=4/1024, ttl=64 (reply in 8)
8	2025-10-23 15:2...	192.168.60.5	10.9.0.5	ICMP	98	Echo (ping) reply id=0x0012, seq=4/1024, ttl=63 (request in 7)
9	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) request id=0x0012, seq=5/1280, ttl=64 (reply in 10)
10	2025-10-23 15:2...	192.168.60.5	10.9.0.5	ICMP	98	Echo (ping) reply id=0x0012, seq=5/1280, ttl=63 (request in 9)
11	2025-10-23 15:2...	f6:6f:bb:cb:71:3c	52:ed:3a:0b:65:2f	ARP	42	Who has 10.9.0.5? Tell 10.9.0.11
12	2025-10-23 15:2...	52:ed:3a:0b:65:2f	f6:6f:bb:cb:71:3c	ARP	42	Who has 10.9.0.11? Tell 10.9.0.5
13	2025-10-23 15:2...	f6:6f:bb:cb:71:3c	52:ed:3a:0b:65:2f	ARP	42	10.9.0.5 is at 52:ed:3a:0b:65:2f
14	2025-10-23 15:2...	f6:6f:bb:cb:71:3c	52:ed:3a:0b:65:2f	ARP	42	10.9.0.11 is at f6:6f:bb:cb:71:3c
15	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) request id=0x0012, seq=6/1536, ttl=64 (reply in 16)
16	2025-10-23 15:2...	192.168.60.5	10.9.0.5	ICMP	98	Echo (ping) reply id=0x0012, seq=6/1536, ttl=63 (request in 15)
17	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) request id=0x0012, seq=7/1792, ttl=64 (reply in 18)
18	2025-10-23 15:2...	192.168.60.5	10.9.0.5	ICMP	98	Echo (ping) reply id=0x0012, seq=7/1792, ttl=63 (request in 17)
19	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) request id=0x0012, seq=8/2048, ttl=64 (reply in 20)
20	2025-10-23 15:2...	192.168.60.5	10.9.0.5	ICMP	98	Echo (ping) reply id=0x0012, seq=8/2048, ttl=63 (request in 19)
21	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) request id=0x0012, seq=9/2304, ttl=64 (reply in 22)
22	2025-10-23 15:2...	192.168.60.5	10.9.0.5	ICMP	98	Echo (ping) reply id=0x0012, seq=9/2304, ttl=63 (request in 21)
23	2025-10-23 15:2...	10.9.0.1	224.0.0.251	MDNS	104	Standard query 0x0000 PTR _nmea._0183._tcp.local, "QM" question PTR _ipp._t...
24	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) request id=0x0012, seq=10/2560, ttl=64 (reply in 25)
25	2025-10-23 15:2...	192.168.60.5	10.9.0.5	ICMP	98	Echo (ping) reply id=0x0012, seq=10/2560, ttl=63 (request in 24)
26	2025-10-23 15:2...	f80:fc12:5cff:fe80:ff02:fb	MDNS	124	Standard query 0x0000 PTR _nmea._0183._tcp.local, "QM" question PTR _ipp._t...	
27	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) request id=0x0012, seq=11/2816, ttl=64 (reply in 28)
28	2025-10-23 15:2...	192.168.60.5	10.9.0.5	ICMP	98	Echo (ping) reply id=0x0012, seq=11/2816, ttl=63 (request in 27)
29	2025-10-23 15:2...	10.9.0.5	192.168.60.5	ICMP	98	Echo (ping) request id=0x0012, seq=12/3072, ttl=64 (reply in 30)
30	2025-10-23 15:2...	192.168.60.5	10.9.0.5	ICMP	98	Echo (ping) reply id=0x0012, seq=12/3072, ttl=63 (request in 29)

▶ Frame 1: 98 bytes on wire (784 bits), 98 bytes captured (784 bits) on interface br-a645e9fb5584, id 0
 ▶ Ethernet II, Src: 52:ed:3a:0b:65:2f (52:ed:3a:0b:65:2f), Dst: f6:6f:bb:cb:71:3c (f6:6f:bb:cb:71:3c)
 ▶ Internet Protocol Version 4, Src: 10.9.0.5, Dst: 192.168.60.5
 ▶ Internet Control Message Protocol

```

0000  f6 6f bb cb 71 3c 52 ed 3a 0b 65 2f 08 00 45 00  ·· q<R· :· e· E·
0010  00 54 b7 a8 40 00 40 01 7c 45 aa 09 00 05 c8 a8  ·T· @ @· |E···
0020  3c 05 08 00 12 8e 00 12 00 01 de 47 fa 68 00 00 <··· ·G· h··
0030  00 00 41 db 0c 00 00 00 00 00 10 11 12 13 14 15  ·A··· ······
0040  16 17 18 19 1a 1b 1c 1d 1e 1f 20 21 22 23 24 25  ······ !#$%
0050  26 27 28 29 2a 2b 2c 2d 2e 2f 30 31 32 33 34 35  &'()*, - ./012345
0060  36 37 67
  
```

br-a645e9fb5584:<live capture in progress> Packets: 30 · Displayed: 30 (100.0%) Profile: Default

Step6: Cleanup

44) Command:

```
iptables -F
iptables -P OUTPUT ACCEPT
iptables -P INPUT ACCEPT
```

```
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -F
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -P OUTPUT ACCEPT
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -P INPUT ACCEPT
seed-router:PES1UG23CS433:PranavHemanth:/
$iptables -t filter -L -n
chain INPUT (policy ACCEPT)
target      prot opt source          destination
chain FORWARD (policy DROP)
target      prot opt source          destination
chain OUTPUT (policy ACCEPT)
target      prot opt source          destination
seed-router:PES1UG23CS433:PranavHemanth:/
$
```

Another way to restore the states of all the tables is to restart the container.

You can do it using the following command (you need to find the container's ID first):

\$ docker restart <container ID> (Here - \$ docker restart dccca3a58e38)

Explanation: Experiment 2:- Without the Second DROP Rule

Here, only the first rule with the limit module was added, and the chain's default policy remained ACCEPT. When pinging 192.168.60.5 from Host A, the connection is fully successful with 0% packet loss, showing that the rate limit is ineffective. The excess packets that would have been blocked by the second rule are instead allowed to pass through.

Purpose of the single rule:

- iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute --limit-burst 5 -j ACCEPT — allows packets within the specified rate, but does not automatically block packets exceeding the limit. Without an explicit DROP rule, any packet beyond the limit is simply allowed by the chain's default ACCEPT policy, defeating the purpose of the rate limit.

Overall Conclusion:

The second DROP rule is essential for enforcing the rate limit. It ensures that any packets exceeding the allowed threshold are explicitly discarded, otherwise the limit module alone cannot restrict traffic effectively.

Task 5: Load Balancing

The iptables are very powerful. In addition to firewalls, it has many other applications. We will not be able to cover all its applications in this lab, but we will be experimenting with one of the applications, load balancing. In this task, we will use it to balance three UDP servers running in the internal network. Let's first start the server on each of the hosts: 192.168.60.5, 192.168.60.6, and 192.168.60.7 (the -k option indicates that the server can receive UDP datagrams from multiple hosts):

```
# nc -luk 8080
```

We can use the statistic module to achieve load balancing.

There are two modes: random and nth. We will conduct experiments using both of them.

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

Step1: We shall implement policies to equally divide the incoming packets between the three interval servers. On the seed-router container:

45)Command:

```
iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth  
--every 3 --packet 0 -j DNAT --to-destination 192.168.60.5:8080  
iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth  
--every 2 --packet 0 -j DNAT --to-destination 192.168.60.6:8080  
iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth  
--every 1 --packet 0 -j DNAT --to-destination 192.168.60.7:8080  
iptables -L -n -v
```

seed-router:

```
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 3 --packet 0 -j DNAT --to-destination 192.168.60.5:8080  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 2 --packet 0 -j DNAT --to-destination 192.168.60.6:8080  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode nth --every 1 --packet 0 -j DNAT --to-destination 192.168.60.7:8080  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -L -n -v  
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)  
  pkts bytes target     prot opt in     out      source               destination  
  
Chain FORWARD (policy ACCEPT 0 packets, 0 bytes)  
  pkts bytes target     prot opt in     out      source               destination  
  
Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)  
  pkts bytes target     prot opt in     out      source               destination  
seed-router:PES1UG23CS433:PranavHemanth:/  
$iptables -t filter -L -n  
Chain INPUT (policy ACCEPT)  
target     prot opt source          destination  
  
Chain FORWARD (policy ACCEPT)  
target     prot opt source          destination  
  
Chain OUTPUT (policy ACCEPT)  
target     prot opt source          destination  
seed-router:PES1UG23CS433:PranavHemanth:/  
$
```

Step2: Start servers on host1 - 192.168.60.5, host2 - 192.168.60.6 and host3 - 192.168.60.7

46)Command: nc -luk 8080

host1 - 192.168.60.5:

```
host1-192.168.60.5:PES1UG23CS433:PranavHemanth:/  
$nc -luk 8080  
[ ]
```

host2 - 192.168.60.6:

```
host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/  
$nc -luk 8080  
[ ]
```

host3 - 192.168.60.7:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
|host3-192.168.60.7:PES1UG23CS433:PranavHemanth:/  
$nc -luk 8080  
|
```

47)Command: Now on hostA - 10.9.0.5

```
nc -u 10.9.0.11 8080  
< enter 3 words, wait 30 seconds before entering the next word>
```

hostA - 10.9.0.5:

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/  
$>nc -u 10.9.0.11 8080  
Pranav  
Hemanth  
PES1UG23CS433  
|
```

host1 - 192.168.60.5:

```
host1-192.168.60.5:PES1UG23CS433:PranavHemanth:/  
$>nc -luk 8080  
Pranav
```

host2 - 192.168.60.6:

```
host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/  
$>nc -luk 8080  
Hemanth
```

host3 - 192.168.60.7:

```
host3-192.168.60.7:PES1UG23CS433:PranavHemanth:/  
$>nc -luk 8080  
PES1UG23CS433  
|
```

Step3: Cleanup

48)Command:

```
iptables -F  
iptables -P OUTPUT ACCEPT  
iptables -P INPUT ACCEPT
```

```
seed-router@PES1UG23CS433:PranavHemanth:/  
$iptables -F  
seed-router@PES1UG23CS433:PranavHemanth:/  
$iptables -P OUTPUT ACCEPT  
seed-router@PES1UG23CS433:PranavHemanth:/  
$iptables -P INPUT ACCEPT  
seed-router@PES1UG23CS433:PranavHemanth:/  
$iptables -t filter -L -n  
Chain INPUT (policy ACCEPT)  
target     prot opt source          destination  
  
Chain FORWARD (policy DROP)  
target     prot opt source          destination  
  
Chain OUTPUT (policy ACCEPT)  
target     prot opt source          destination  
seed-router@PES1UG23CS433:PranavHemanth:/  
$
```

Explanation: Experiment 1: Round-Robin Mode (Nth Packet)

In this experiment, the router is configured to distribute incoming UDP traffic on port 8080 evenly across three internal servers (192.168.60.5, 192.168.60.6, 192.168.60.7) using the statistic module in nth mode. Three DNAT rules are added to the PREROUTING chain of the nat table. The first rule redirects every third packet to Host1, the second redirects every second packet to Host2, and the final rule acts as a catch-all for the remaining packets, sending them to Host3.

Testing this setup from HostA, each UDP packet sent to the router is correctly distributed: the first packet reaches Host1, the second reaches Host2, and the third reaches Host3. This confirms that the router successfully implements round-robin load balancing, ensuring that the packets are evenly split across all servers in a predictable sequence.

Purpose of the rules:

- --every 3 --to-destination 192.168.60.5:8080 — sends every third packet to Host1.
- --every 2 --to-destination 192.168.60.6:8080 — sends every second packet to Host2.
- --every 1 --to-destination 192.168.60.7:8080 — catch-all rule for remaining packets to Host3.

Another way to restore the states of all the tables is to restart the container.

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

You can do it using the following command (you need to find the container's ID first):
\$ docker restart <container ID> (Here - \$ docker restart dccca3a58e38)

Step4: Using the random mode. The following rule will select a matching packet with the probability P . You need to replace P with a probability number.

Command:

```
iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.3333 -j DNAT --to-destination 192.168.60.5:8080
```

```
iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.5 -j DNAT --to-destination 192.168.60.6:8080
```

```
iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 1 -j DNAT --to-destination 192.168.60.6:8080
```

```
iptables -L -n -v
```

seed-router:

```
seed-router:PES1UG23CS433:PranavHemanth:/
Siptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.3333 -j DNAT --to-destination 192.168.60.5:8080
seed-router:PES1UG23CS433:PranavHemanth:/
Siptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.5 -j DNAT --to-destination 192.168.60.6:8080
seed-router:PES1UG23CS433:PranavHemanth:/
Siptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 1 -j DNAT --to-destination 192.168.60.6:8080
seed-router:PES1UG23CS433:PranavHemanth:/
Siptables -L -n -v
Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
  pkts bytes target     prot opt in     out    source          destination
Chain FORWARD (policy ACCEPT 0 packets, 0 bytes)
  pkts bytes target     prot opt in     out    source          destination
Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
  pkts bytes target     prot opt in     out    source          destination
seed-router:PES1UG23CS433:PranavHemanth:/
Siptables -t filter -L -n
Chain INPUT (policy ACCEPT)
  target     prot opt source          destination
Chain FORWARD (policy ACCEPT)
  target     prot opt source          destination
Chain OUTPUT (policy ACCEPT)
  target     prot opt source          destination
seed-router:PES1UG23CS433:PranavHemanth:/
$
```

Step5: Start servers on host1 - 192.168.60.5, host2 - 192.168.60.6 and host3 - 192.168.60.7

49)Command: nc -luk 8080

host1 - 192.168.60.5:

```
host1-192.168.60.5:PES1UG23CS433:PranavHemanth:/
$nc -luk 8080

```

host2 - 192.168.60.6:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/
$nc -luk 8080
```

host3 - 192.168.60.7:

```
host3-192.168.60.7:PES1UG23CS433:PranavHemanth:/
:$nc -luk 8080
```

50)Command: Now on hostA - 10.9.0.5

```
nc -u 10.9.0.11 8080
```

< enter 3 words, wait 30 seconds before entering the next word>

All windows:

```
host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/
$nc -luk 8080
Hello 2
host3-192.168.60.7:PES1UG23CS433:PranavHemanth:/
:$nc -luk 8080
Hello 3
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/
$nc -u 10.9.0.11 8080
Hello 1
Hello 2
Hello 3
```

hostA - 10.9.0.5:

```
hostA-10.9.0.5:PES1UG23CS433:PranavHemanth:/
$nc -u 10.9.0.11 8080
Hello 1
Hello 2
Hello 3
```

host1 - 192.168.60.5:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

```
host1-192.168.60.5:PES1UG23CS433:PranavHemanth:/  
$nc -luk 8080  
Hello 1  
[]
```

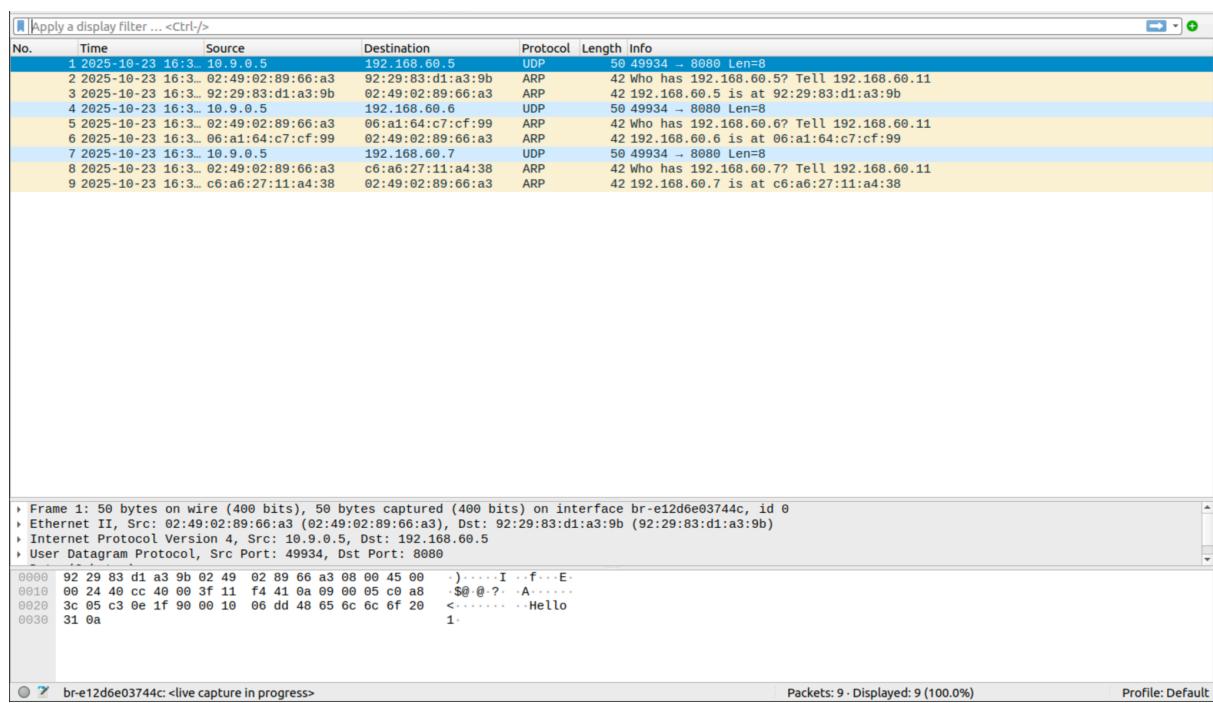
host2 - 192.168.60.6:

```
host2-192.168.60.6:PES1UG23CS433:PranavHemanth:/  
$nc -luk 8080  
Hello 2  
[]
```

host3 - 192.168.60.7:

```
host3-192.168.60.7:PES1UG23CS433:PranavHemanth:/  
$nc -luk 8080  
Hello 3  
[]
```

Wireshark:



The above is not a good illustration of the random mode. Hence the below experiment to try packet probability:

All windows:

Aug -Dec 2025 Assignment SUBMISSION_UE23CS343AB6

The image displays three terminal windows showing network traffic distribution between three hosts (Host1, Host2, and Host3) using the nmap tool.

- Host1 (192.168.60.5):** Shows traffic to port 8080. It receives 1 packet from HostA (10.9.0.5) and 1 packet from Host2 (192.168.60.6). Both packets contain "Hello 1".
- Host2 (192.168.60.6):** Shows traffic to port 8080. It receives 2 packets from HostA (10.9.0.5) and 1 packet from Host1 (192.168.60.5). One packet from HostA contains "Hello 2" and one from Host1 contains "Hello 2".
- Host3 (192.168.60.7):** Shows traffic to port 8080. It receives 3 packets from HostA (10.9.0.5). All three packets contain "Hello 3".

We are not able to view easily the difference since the packets hit probability incidentally in a way consistent to the non random mode

Explanation: Experiment 2: Random Mode (Probability)

In this experiment, the router is configured to distribute UDP traffic using the random mode of the statistic module. Three DNAT rules are added to the PREROUTING chain with cascading probabilities: 0.3333 for Host1, 0.5 for Host2, and 1 for Host3. Each incoming packet is evaluated probabilistically: it may hit one of the first two rules with the defined chance or be caught by the final catch-all rule.

Testing this setup from HostA, the three packets sent to the router are distributed across the three servers according to their probabilities. The first packet hits the 33.33% chance and is sent to Host1, the second misses the first but hits the 50% chance of the second rule and is sent to Host2, and the third misses both probabilistic rules and is sent to Host3. This demonstrates that random-mode load balancing achieves an approximate even distribution of traffic in a non-deterministic manner.

Purpose of the rules:

- --probability 0.3333 --to-destination 192.168.60.5:8080 — sends packets to Host1 with ~1/3 probability.
- --probability 0.5 --to-destination 192.168.60.6:8080 — sends remaining packets to Host2 with ~1/3 probability.
- --probability 1 --to-destination 192.168.60.7:8080 — catch-all sends all leftover packets to Host3.

Overall Conclusion:

Round-robin mode provides predictable, evenly spaced packet distribution, while random mode provides probabilistic load balancing, which approximates equal distribution but with some variability. Both methods allow iptables to function as a simple UDP load balancer, redirecting incoming traffic across multiple internal servers.